

AI for video games

Fabrice Lamarche
ESIR

Schedule

- ◆ Motion planning
 - Deterministic methods
 - Potential fields
 - Probabilistic methods
- ◆ Steering behaviors
 - ◆ Path following
 - ◆ Collision avoidance
 - ◆ Groups
- ◆ Search algorithms
 - Dijkstra, Wavefront, A*, IDA*
- ◆ AI
 - Rule based systems
 - State machines
 - Hierarchical Task networks

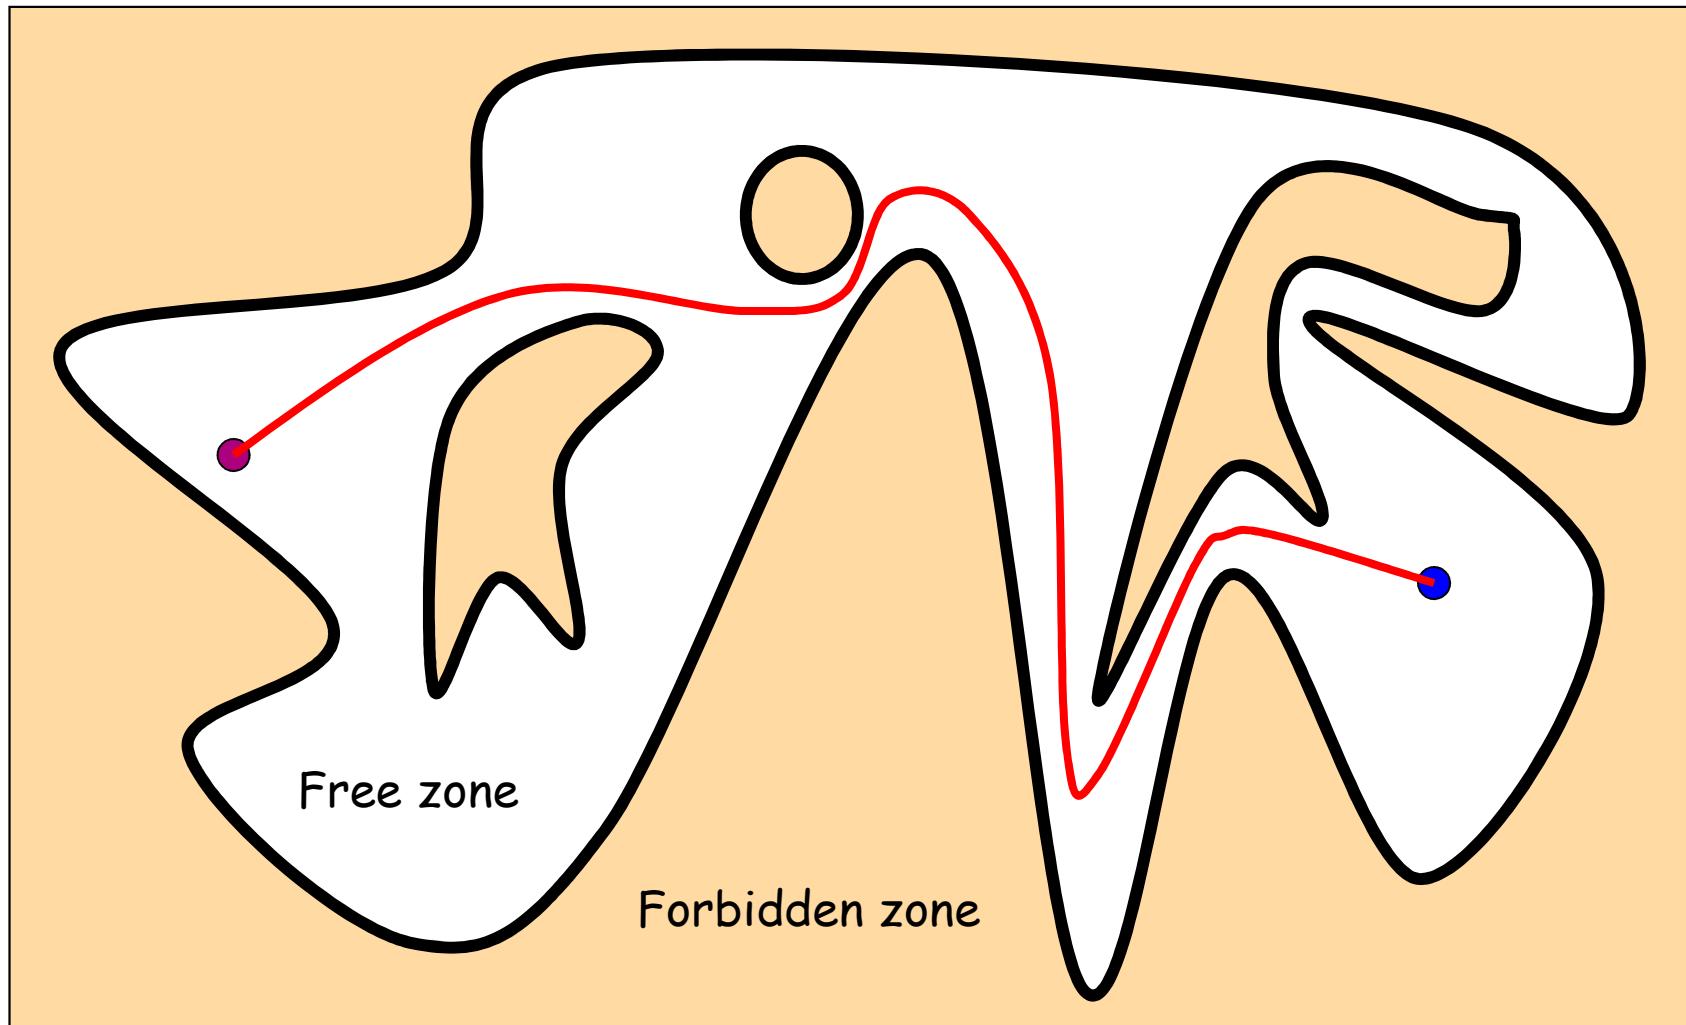
Some ressources

- ◆ <http://www.aigamedev.com>
- ◆ AI Game programming Wisdom. ISBN: 1584500778.
- ◆ Steven M. LaValle (2006). *Planning Algorithms*. Cambridge University Press.
 - <http://planning.cs.uiuc.edu/>
- ◆ Funge (2004). *Artificial Intelligence for Computer Games: An Introduction*. A K Peters. ISBN 1-56881-208-6.
- ◆ Bourg; Seemann (2004). *AI for Game Developers*. O'Reilly & Associates. ISBN 0-596-00555-5.
- ◆ Millington (2005). *Artificial Intelligence for Games*. Morgan Kaufmann. ISBN 0-12-497782-0.

Motion planning

The main question:

Are two points connected with a path?

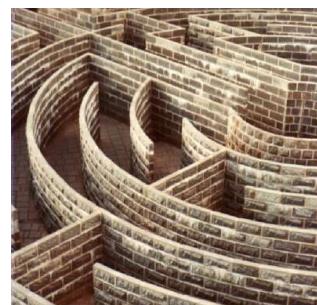
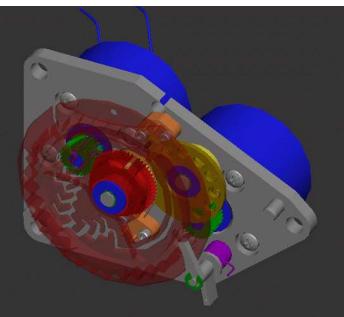


Objectives

- ◆ Computation of movement strategies
 - Geometric paths
 - Trajectories
- ◆ Fulfill high level goals
 - Reach a target without collision with obstacles
 - Product assembly
 - ...

Applications

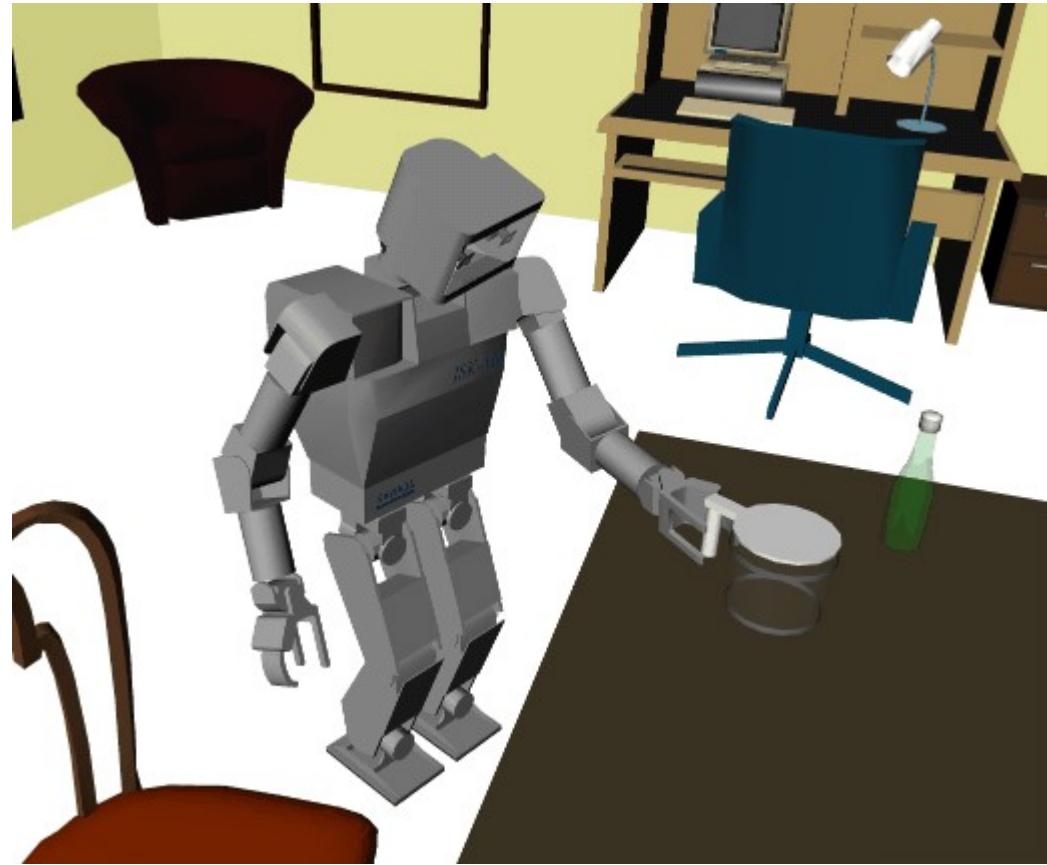
- ◆ Robot programming (painting, assembly)
- ◆ Autonomous vehicles (transport, planet exploration, military applications)
- ◆ Computer animation (video games, films)
- ◆ Surgery (implants)
- ◆ Molecular biology (medicine)



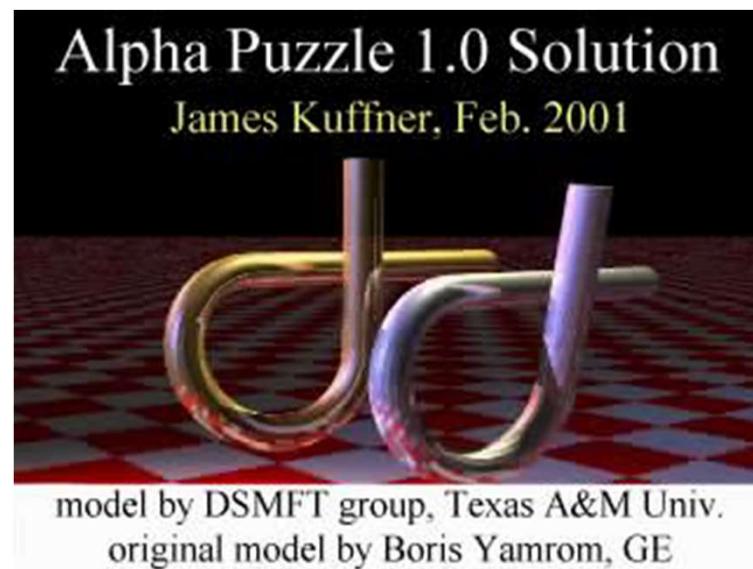
Examples



Examples



Examples



The piano mover's problem

- ◆ Given a piano and an environment populated with obstacles, can we move the piano from one position / orientation (called configuration) to another without colliding obstacles?



Définition

- ◆ Given:
 - A mechanical system:
 - » Geometry
 - » Mobility
 - An environment
 - Two configurations: initial (source) and final (target)
- ◆ We want to compute a path, without collision, joining the initial and final configurations.

Definition of a configuration

- ◆ A configuration gives the positions of all points of a mechanical system in a given coordinate system.
- ◆ A configuration is a vector of generalized coordinates.

$$\begin{bmatrix} q_1 \\ q_2 \\ q_3 \\ q_4 \\ q_5 \\ q_6 \end{bmatrix}$$

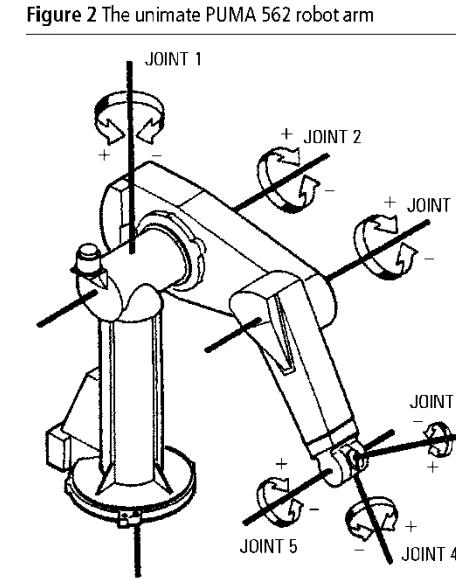
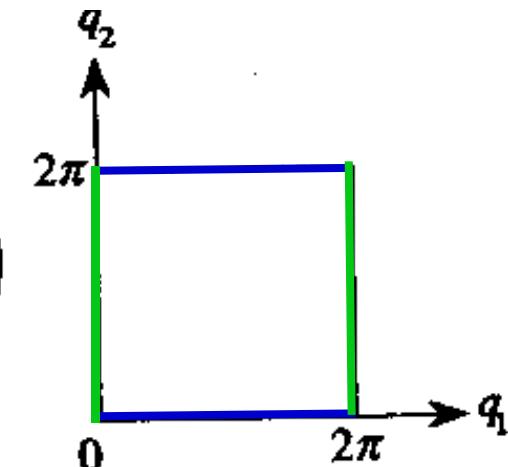
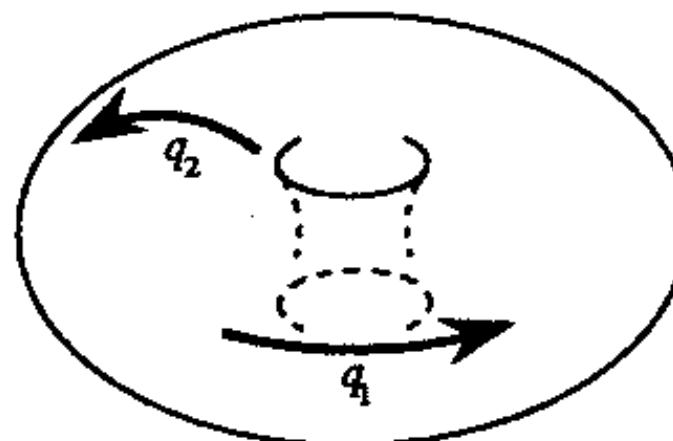
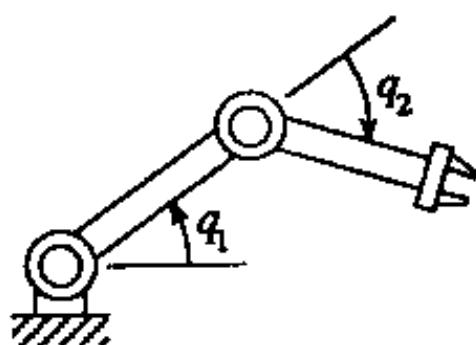


Figure 2 The unimate PUMA 562 robot arm

The configuration space (C-space)

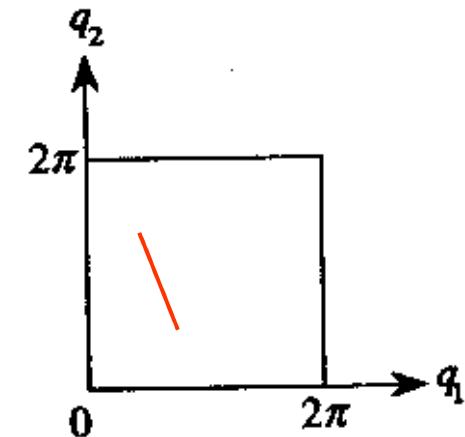
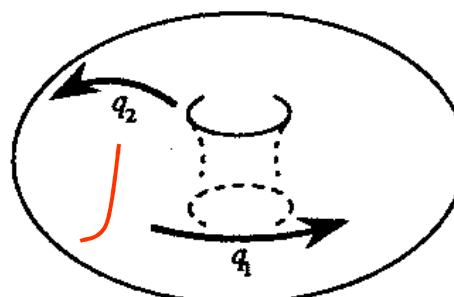
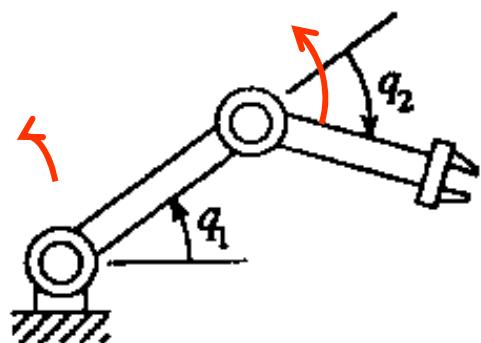
- ◆ Contains all possible configurations of a given mechanical system
- ◆ Formalism described by [Lozano-Perez '79] (C-space)



Path vs trajectory

- ◆ Path: a continuous sequence of configurations
- ◆ Trajectory: temporal parameterization of a path

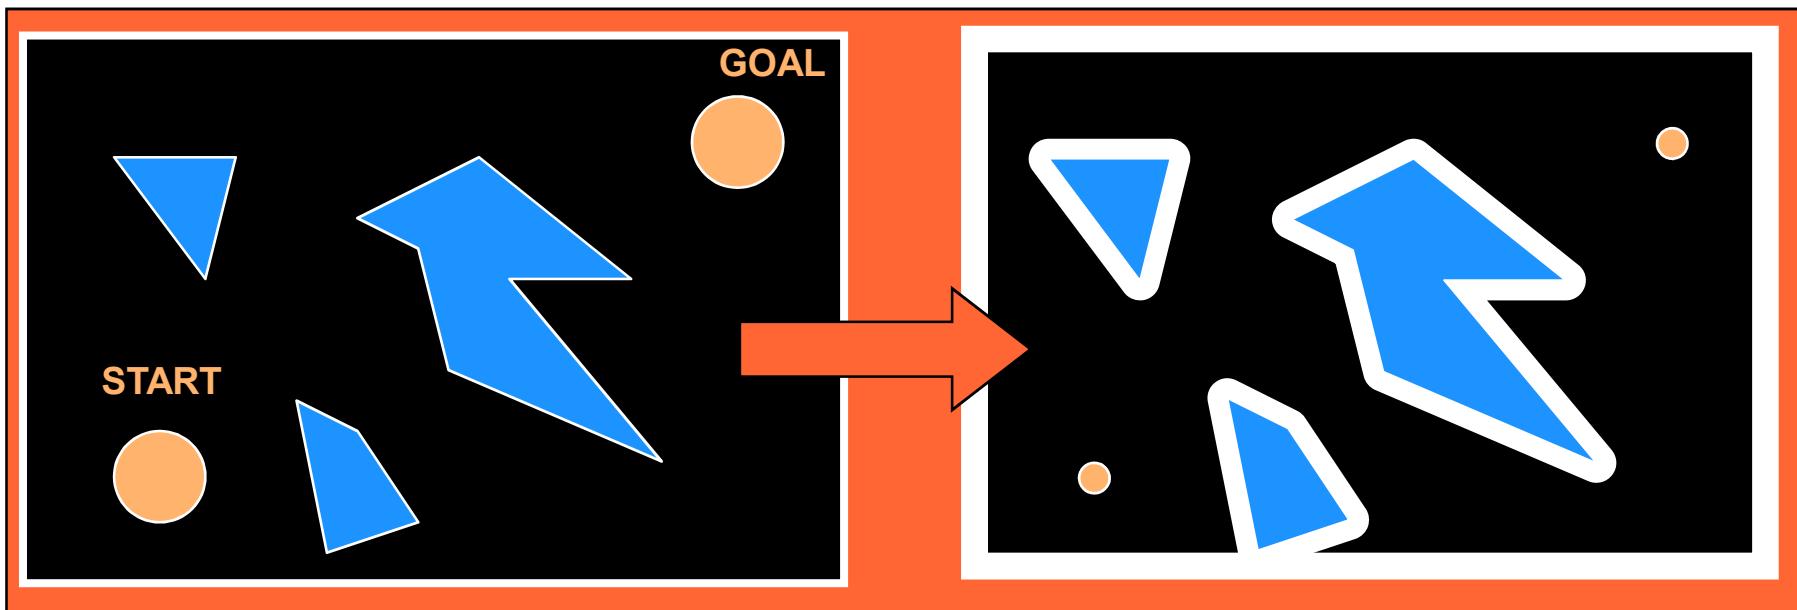
$$q(t) = \begin{bmatrix} q_1(t) \\ q_2(t) \\ \vdots \\ q_n(t) \end{bmatrix}, \quad t \in [0, T]$$



The free space (C-free)

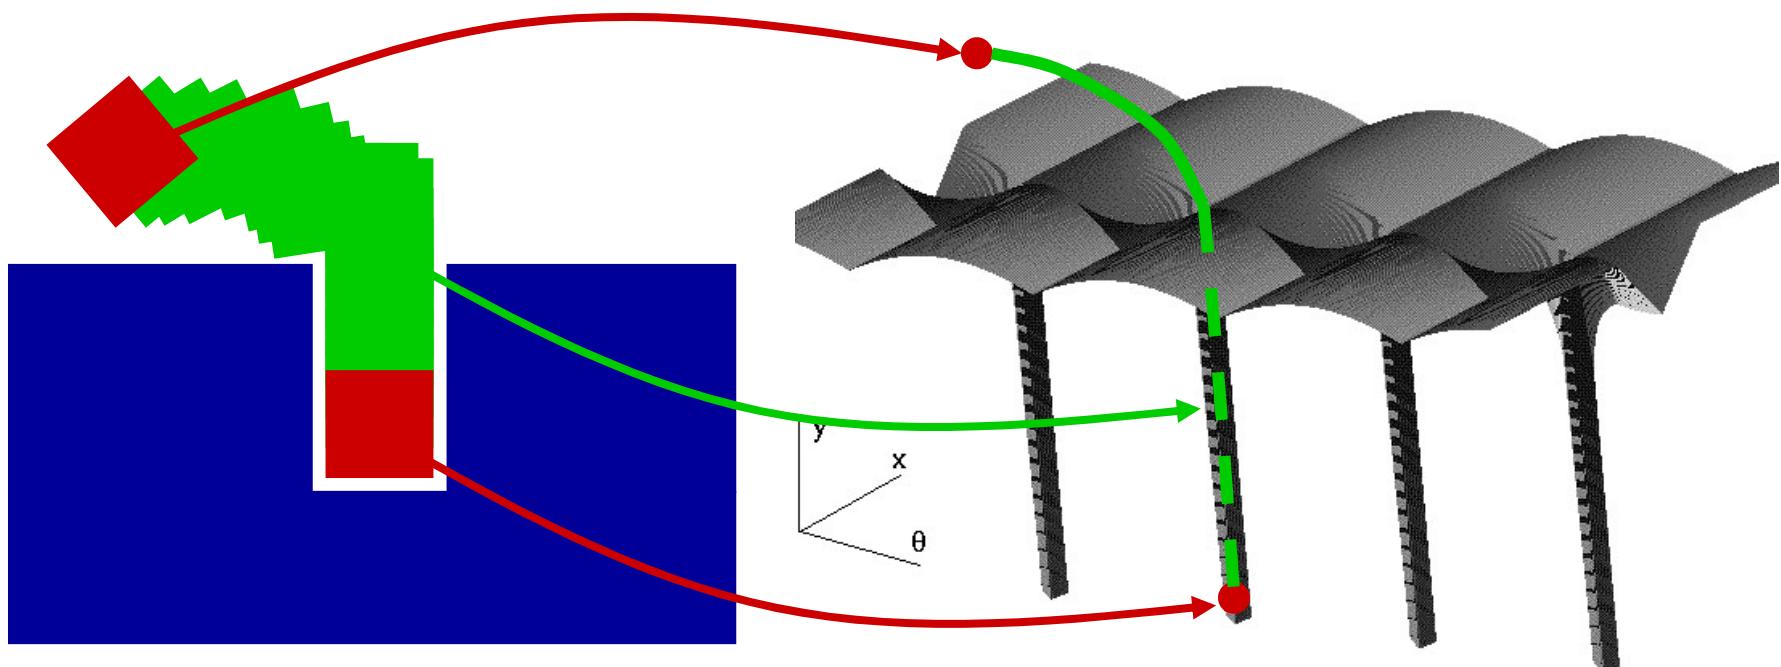
Example: 2D navigation

$$q = \begin{bmatrix} x \\ y \end{bmatrix}$$



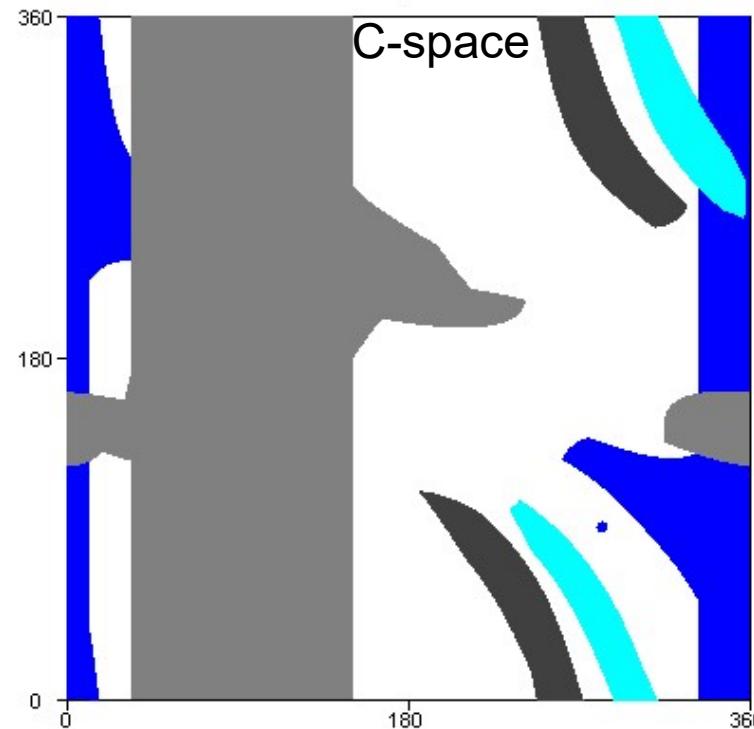
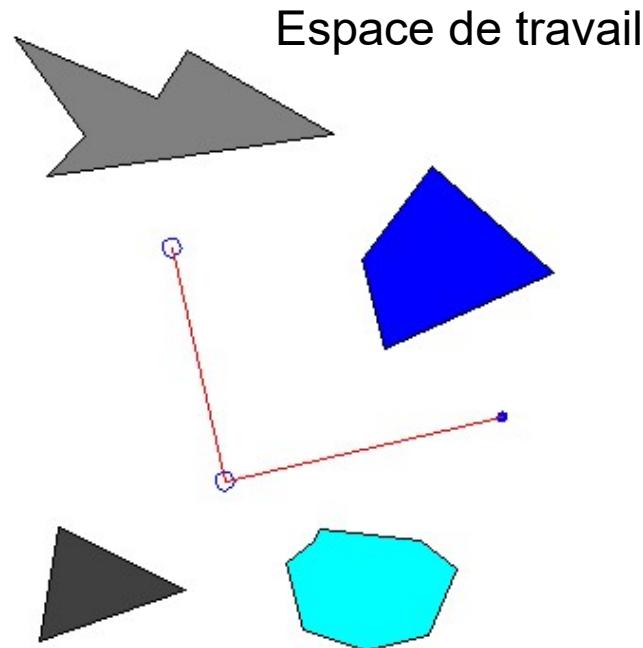
C-space representation

$$C_{space} = C_{free} + C_{obst}$$



C-space representation

$$C_{space} = C_{free} + C_{obst}$$



Images created by C-space Java applet : <http://ford.ieor.berkeley.edu/cspace/>

Main approaches

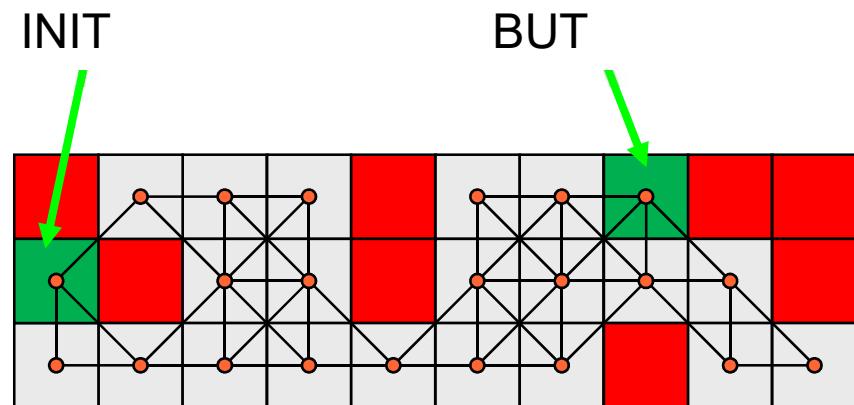
- ◆ Deterministic search
 - Approximate: grids
 - Exact: Cell decompositions
- ◆ Navigation functions: potential fields
- ◆ Probabilistic sampling: exploring high dimensional spaces
 - PRM (Probabilistic RoadMaps)
 - RRT (Rapidly exploring Random Trees)

Main approaches

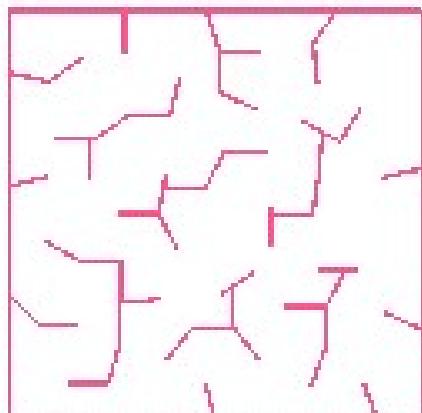
- ◆ Deterministic search
 - Approximate: grids
 - Exact: Cell decompositions
- ◆ Navigation functions: potential fields
- ◆ Probabilistic sampling: exploring high dimensional spaces
 - PRM (Probabilistic RoadMaps)
 - RRT (Rapidly exploring Random Trees)

Approximate methods: Grids

- ◆ Approximate representation of the C-space (C-free and C-obstacle)
- ◆ Each cell is marked free or obstacle



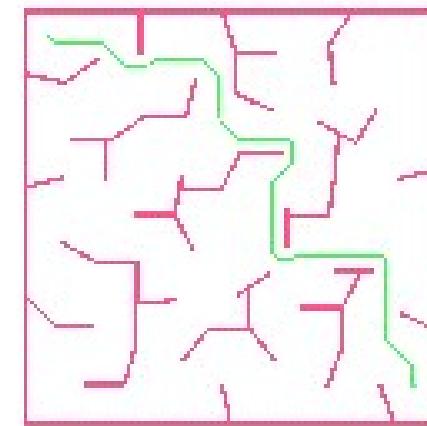
PLANNING: (Dijkstra, A*...)



Projected map



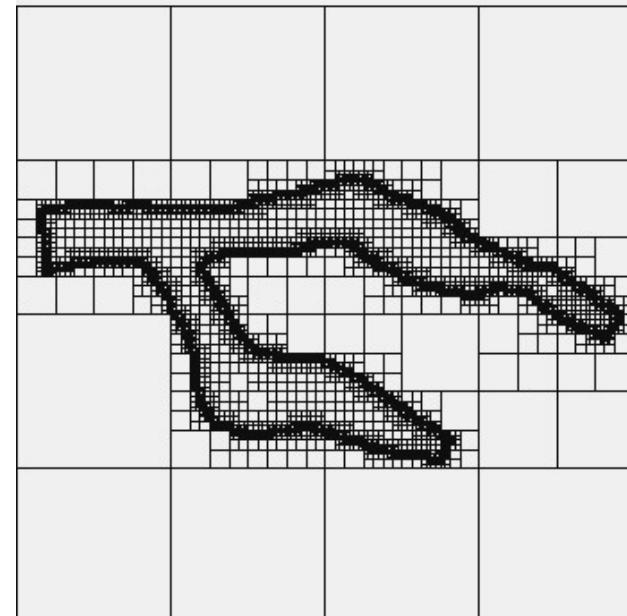
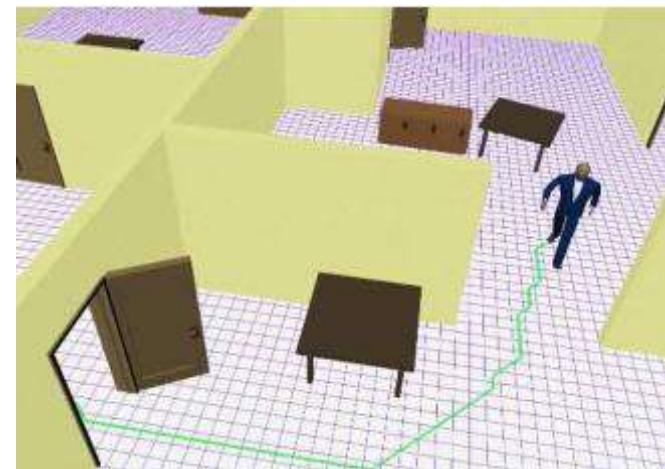
Extended obstacles



Planned path

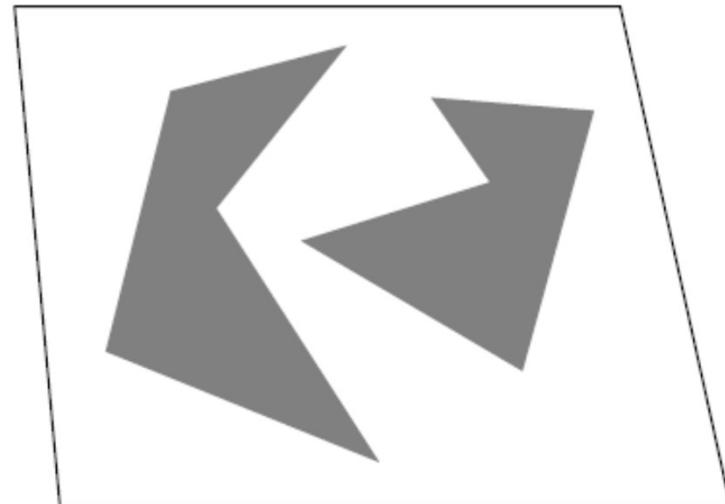
Grids

- ◆ Pros:
 - Easy to implement
 - Fast to compute
- ◆ Cons:
 - Lack of precision
 - Huge search time in high dimensional problems
- ◆ Can be improved
 - Quadtree / Octree
 - More generally: hierarchical representations

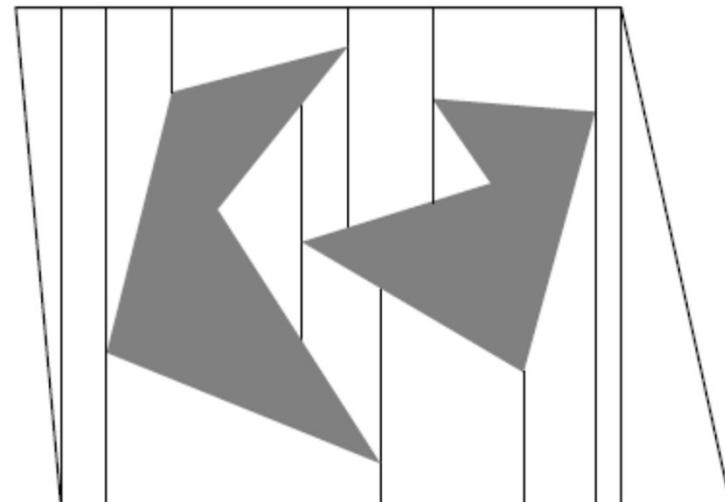


Exact methods: cell decomposition

- ◆ Principle:
 - Decompose C-free into cells of simple geometric shapes (triangles, quads...)

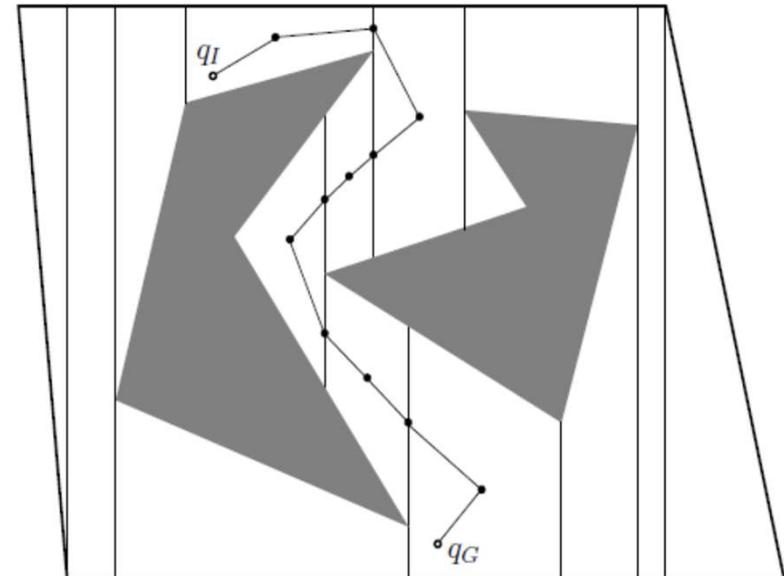
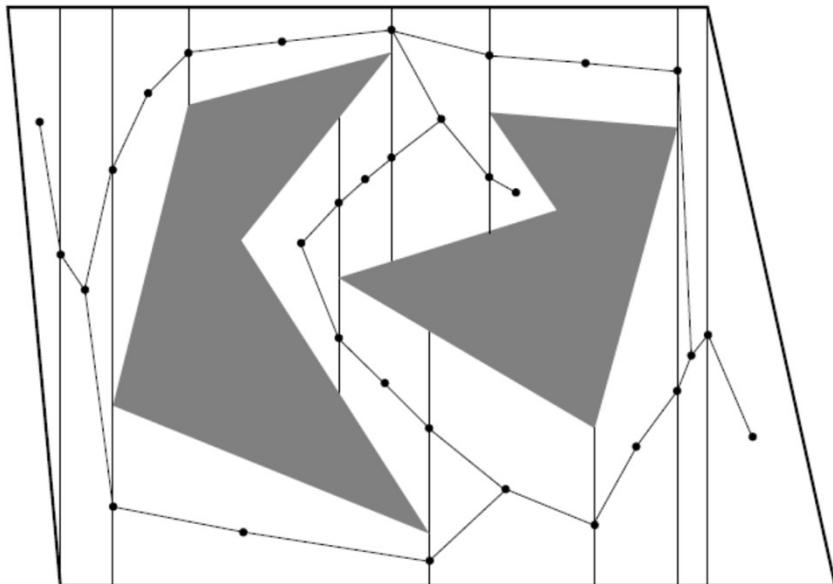


- ◆ Example:
 - Vertical decomposition



Exact cell decomposition

- ◆ A roadmap, capturing C-free topology, is deduced from the decomposition
- ◆ Initial and final configurations are connected to the graph and a path can be planned



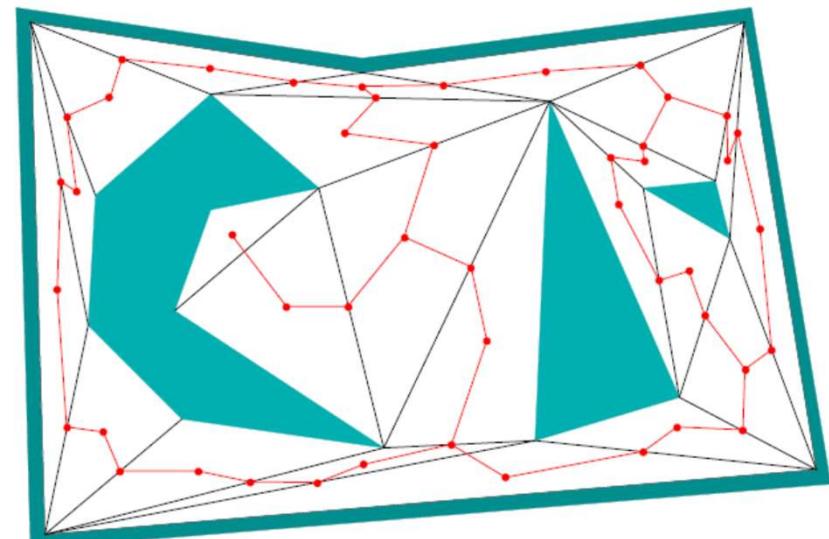
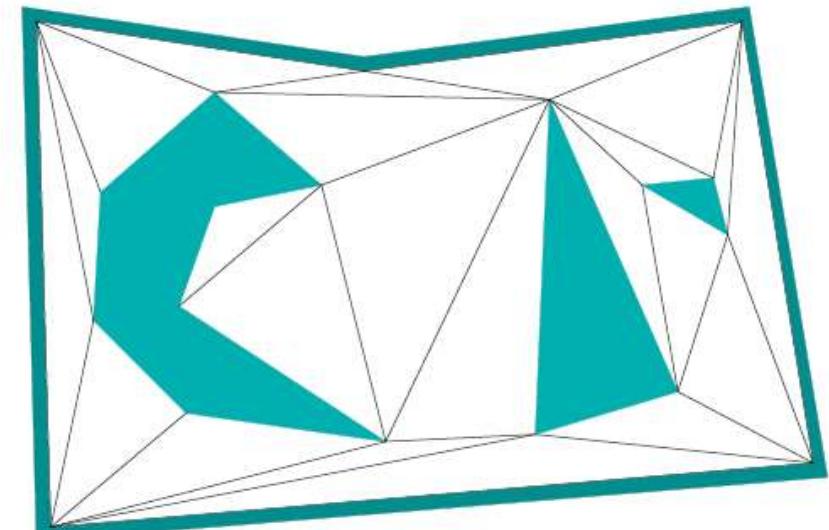
Exact cell decomposition

- ◆ Constrained Delaunay triangulation

- Computes a partition of the plane with triangles
- Obstacles borders are constrained segments

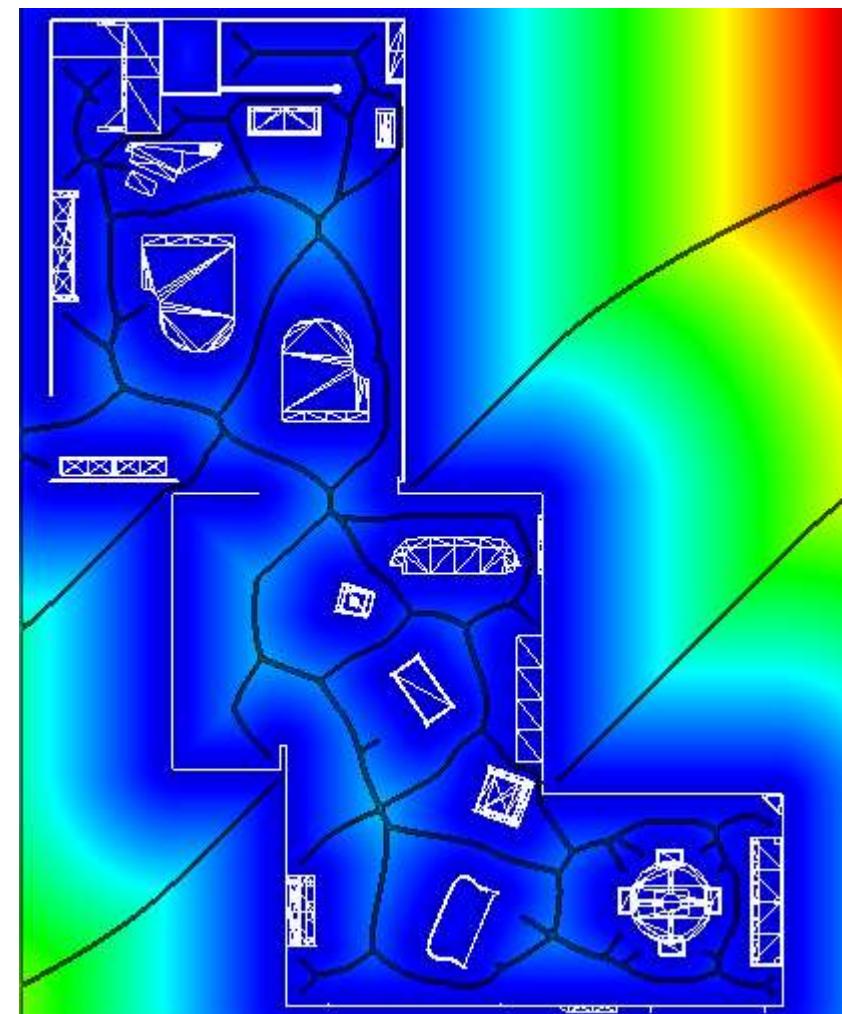
- ◆ Some improvements

- Computation of bottlenecks



Exact methods: Voronoi diagrams

- ◆ Generalized Voronoi diagrams
- ◆ Useful in robotics
 - Computation of safest path
(far from obstacles)

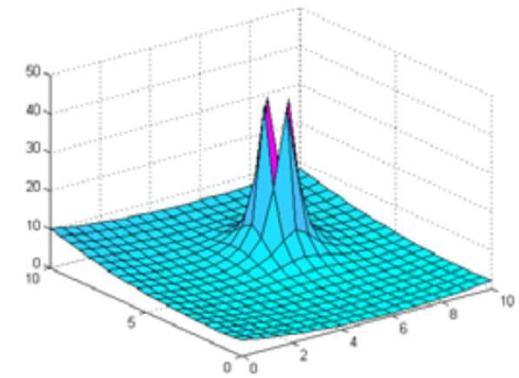
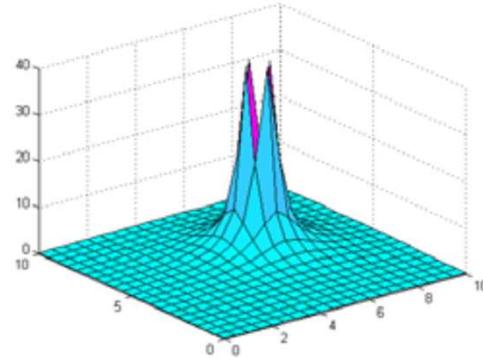
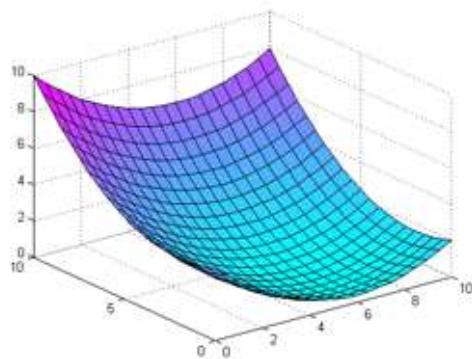


Main approaches

- ◆ Deterministic search
 - Approximate: grids
 - Exact: Cell decompositions
- ◆ **Navigation functions: potential fields**
- ◆ Probabilistic sampling: exploring high dimensional spaces
 - PRM (Probabilistic RoadMaps)
 - RRT (Rapidly exploring Random Trees)

Potential fields

- ◆ Function defined on C-free
- ◆ Ideally:
 - Push mobile far from obstacles
 - Attract mobile toward the goal

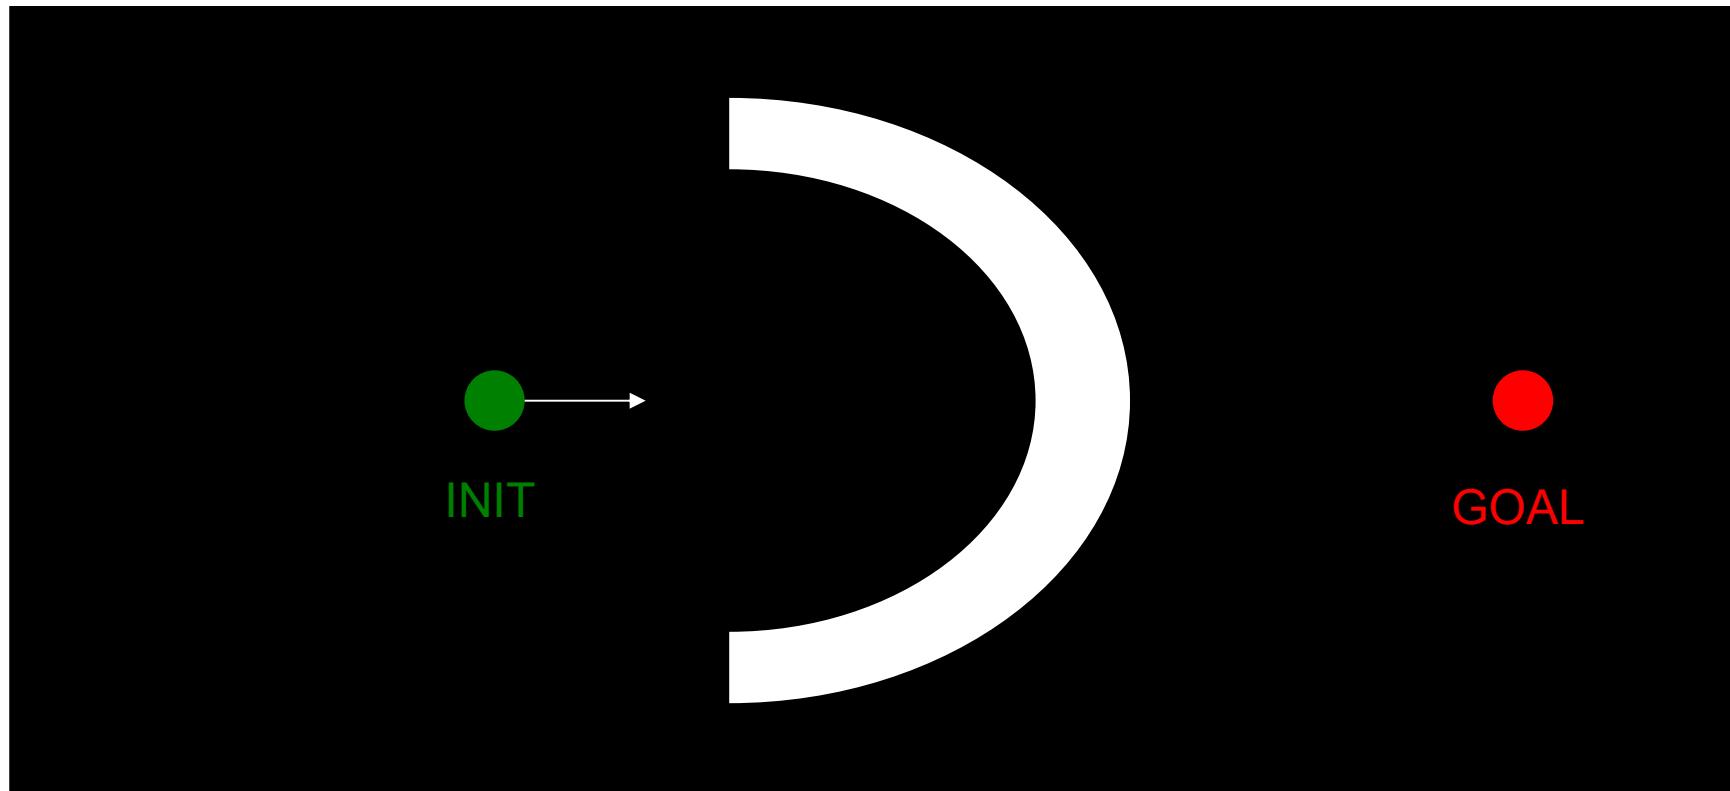


Potential fields

- ◆ An attractive solution for planning problems
 - Paths are not planned but generated in real time
 - Planning and control are merged in a function
 - Paths are smoothed
 - Potential fields can be updated dynamically
 - To reduce computation, only local obstacles can be taken into account

Potential fields

- ◆ But: local minima problem



Potential fields

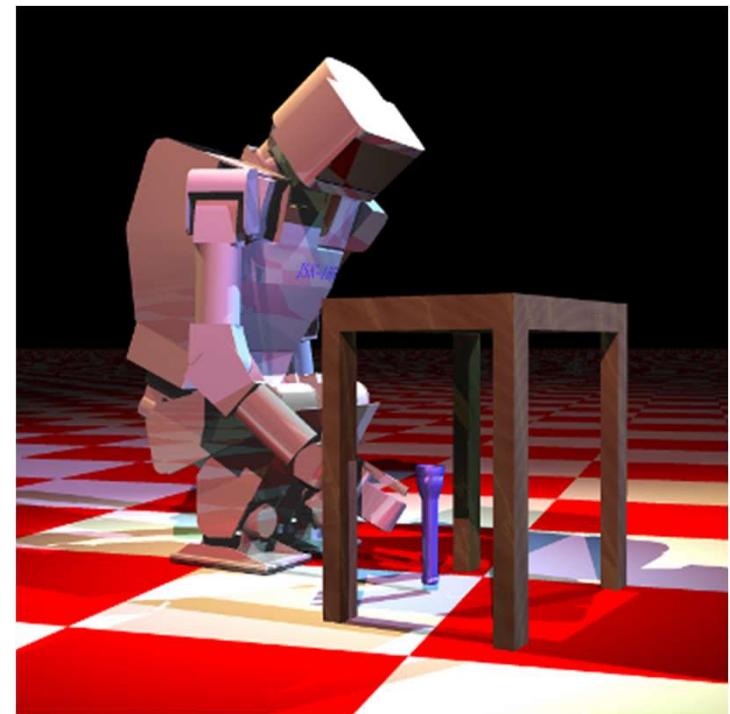
- ◆ Solutions to the local minima problem:
 - Backward chaining
 - Random walk
 - Use a roadmap
- ◆ Pre-computation:
 - Potential field definition
 - Using a regular grid to store potentials
- ◆ Planning:
 - Search in the grid by using potential fields as a heuristic

Main approaches

- ◆ Deterministic search
 - Approximate: grids
 - Exact: Cell decompositions
- ◆ Navigation functions: potential fields
- ◆ **Probabilistic sampling: exploring high dimensional spaces**
 - **PRM (Probabilistic RoadMaps)**
 - **RRT (Rapidly exploring Random Trees)**

Planning in high dimensional space

- ◆ Ideal: a complete planner
 - Guarantees to find a solution in a finite time if it exists
 - Indicates that no solution exists in a finite time
 - Remark: completeness implies the computation of at least one path among all possible paths
- ◆ Problem:
 - PSPACE complete
- ◆ Remark:
 - Exact deterministic methods are complete
 - Approximate deterministic methods are complete in representation



Planning in high dimensional space

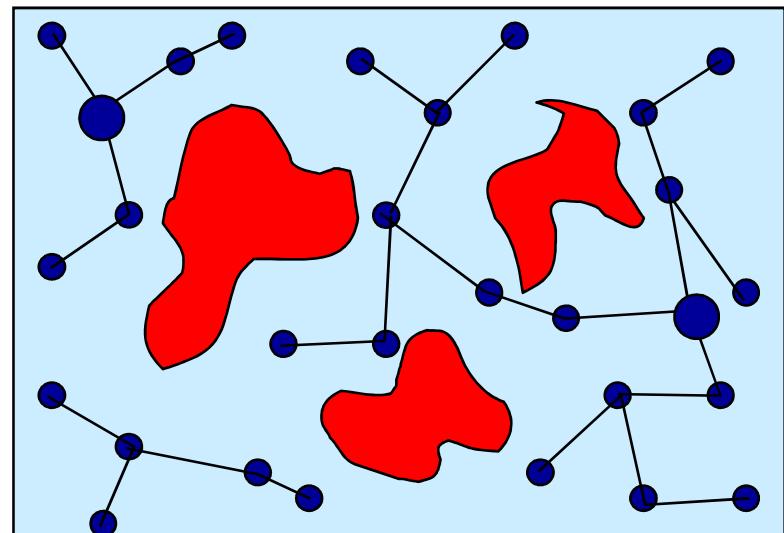
- ◆ Solutions
 - Decrease the number of dimensions
 - Limit the number of possible solutions
 - » Add constraints
 - » Reduce search volume
 - Sacrifice optimality
 - Sacrifice completeness

Planning in high dimensional space

- ◆ Key idea: random exploration of the free space to capture the result
- ◆ Easy way of capturing C-free topology and faster than an exhaustive exploration
- ◆ The problem?
 - Sacrifice of completeness and optimality
 - *Compromise between quality and performances*

PRM : Probabilistic RoadMaps

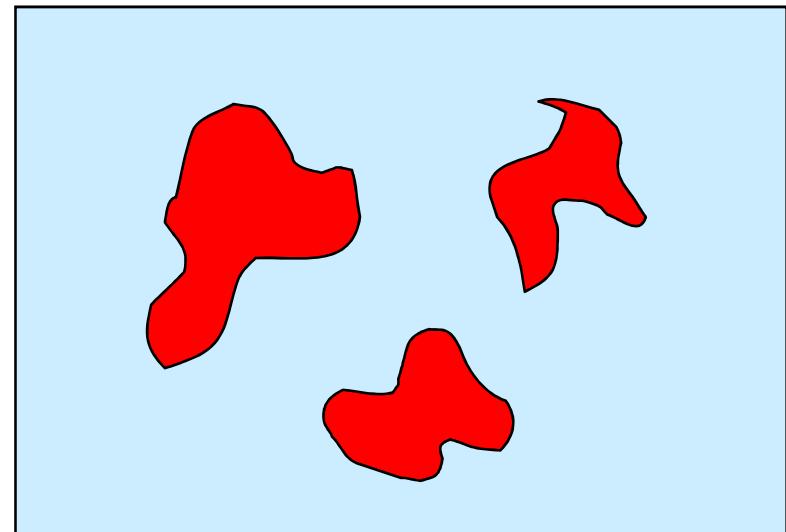
- ◆ Three components
 - Collision detection
 - Local method
 - Sampling
- ◆ Two steps
 - Exploration phase
 - Query phase
- ◆ Key idea: random exploration of C-free to capture its topology in a roadmap
- ◆ Complete in infinite time:
probabilistic completeness



The exploration phase

Computation of the roadmap

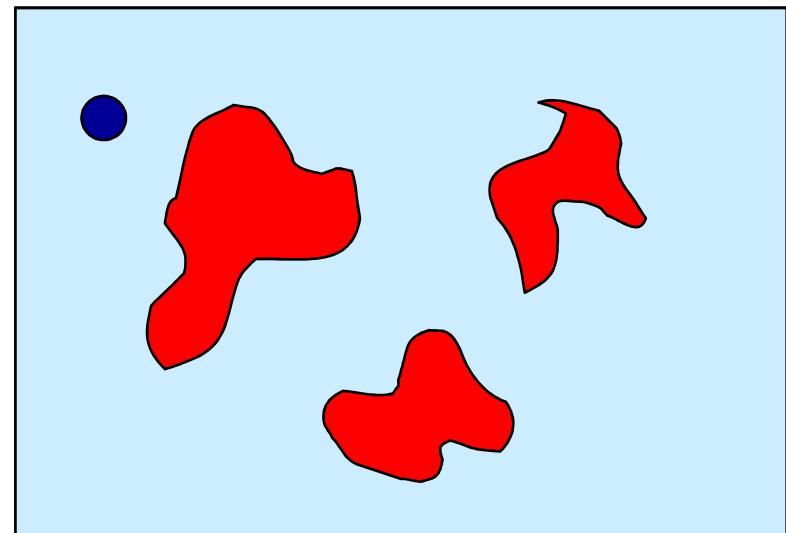
- ◆ Iterative algorithm
 - 1. Computation of random configurations
 - Collision detection
 - 2. Linking configuration
 - Local method
 - » Interpolation between configurations
 - » Collision detection
 - 3. Return to 1



The exploration phase

Computation of the roadmap

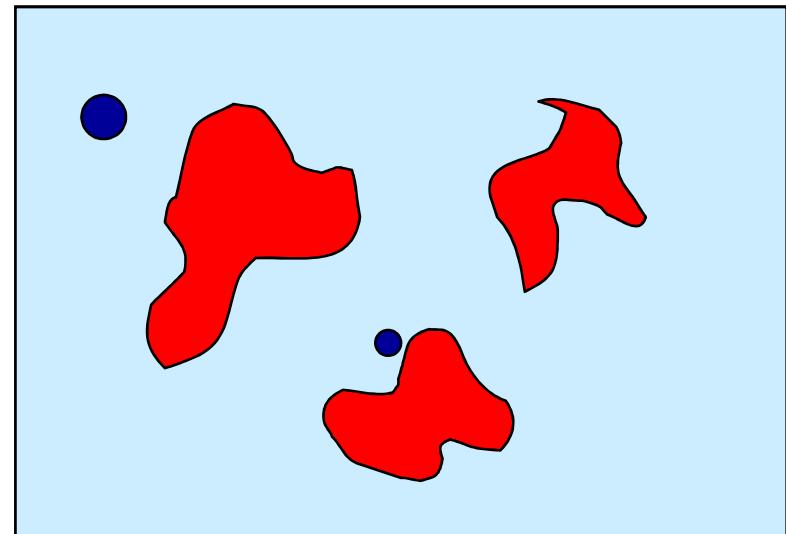
- ◆ Iterative algorithm
 - 1. Computation of random configurations
 - Collision detection
 - 2. Likig configuration
 - Local method
 - » Interpolation between configurations
 - » Collision detection
 - 3. Return to 1



The exploration phase

Computation of the roadmap

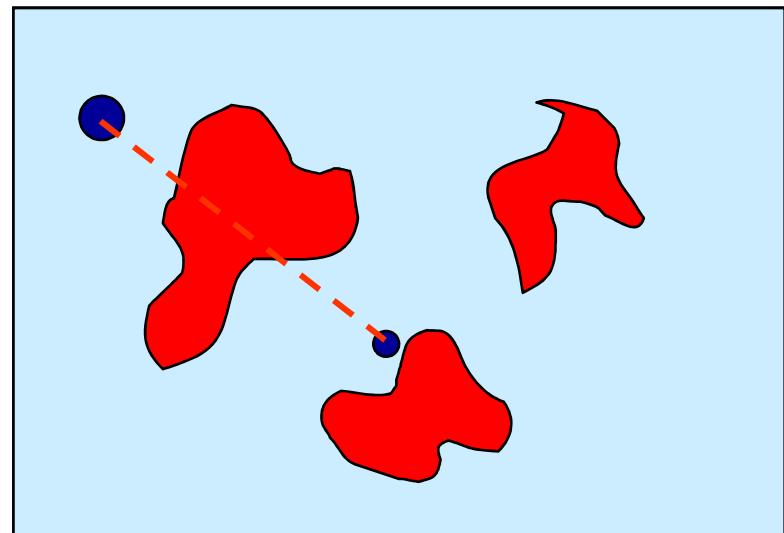
- ◆ Iterative algorithm
 - 1. Computation of random configurations
 - Collision detection
 - 2. Linking configuration
 - Local method
 - » Interpolation between configurations
 - » Collision detection
 - 3. Return to 1



The exploration phase

Computation of the roadmap

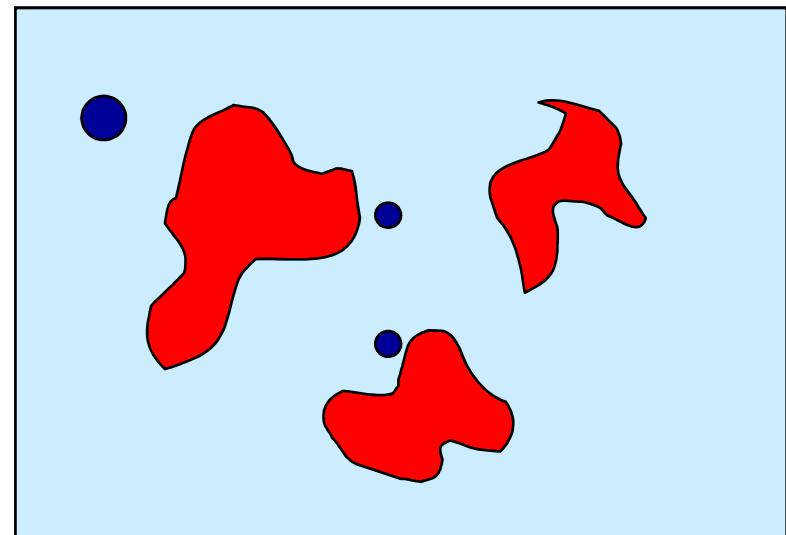
- ◆ Iterative algorithm
 - 1. Computation of random configurations
 - Collision detection
 - 2. Linking configuration
 - Local method
 - » Interpolation between configurations
 - » Collision detection
 - 3. Return to 1



The exploration phase

Computation of the roadmap

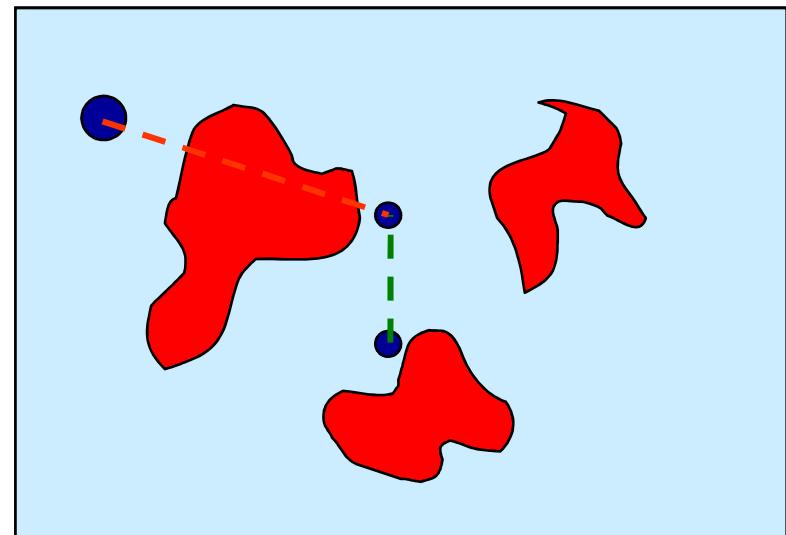
- ◆ Iterative algorithm
 - 1. Computation of random configurations
 - Collision detection
 - 2. Linking configuration
 - Local method
 - » Interpolation between configurations
 - » Collision detection
 - 3. Return to 1



The exploration phase

Computation of the roadmap

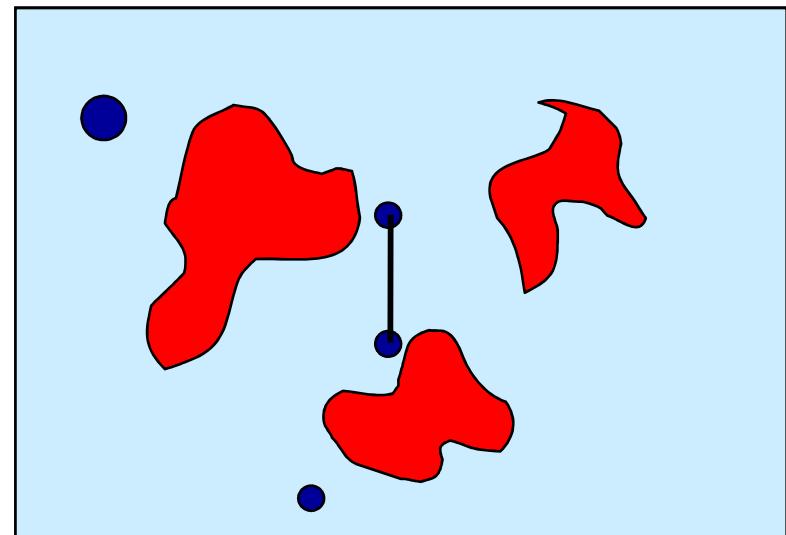
- ◆ Iterative algorithm
 - 1. Computation of random configurations
 - Collision detection
 - 2. Linking configuration
 - Local method
 - » Interpolation between configurations
 - » Collision detection
 - 3. Return to 1



The exploration phase

Computation of the roadmap

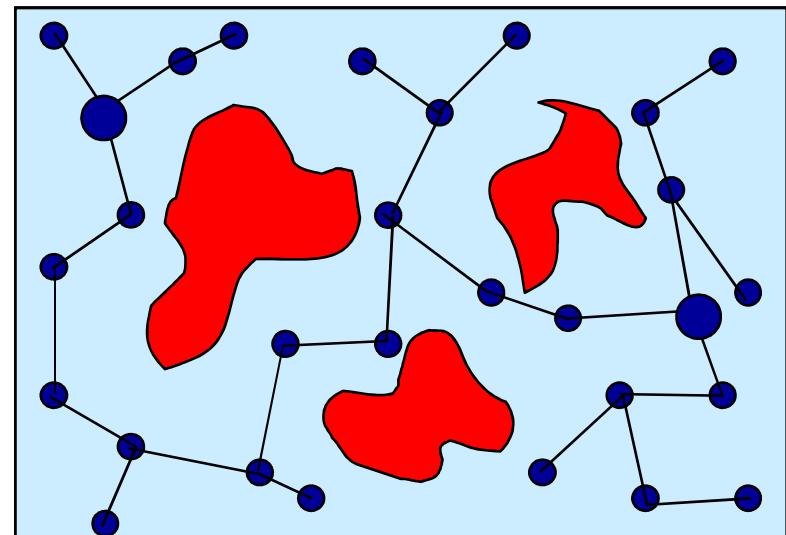
- ◆ Iterative algorithm
 - 1. Computation of random configurations
 - Collision detection
 - 2. Linking configuration
 - Local method
 - » Interpolation between configurations
 - » Collision detection
 - 3. Return to 1



The exploration phase

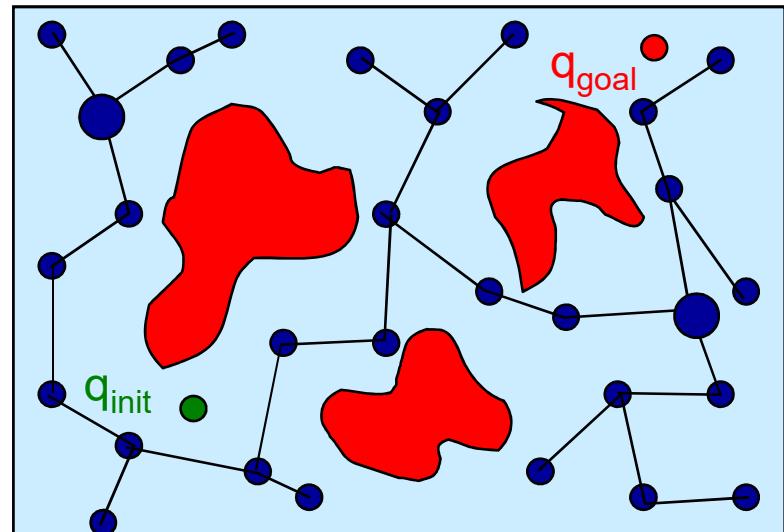
Computation of the roadmap

- ◆ Iterative algorithm
1. Computation of random configurations
 - Collision detection
 2. Linking configuration
 - Local method
 - » Interpolation between configurations
 - » Collision detection
 3. Return to 1



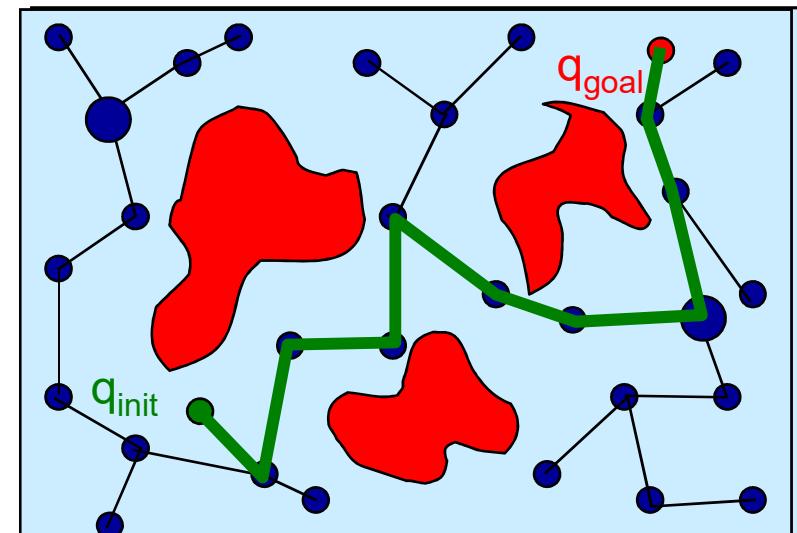
The query phase

- ◆ The PRM is used to solve queries
 - 1. Connect initial and final configurations to the roadmap
 - 2. If nodes belong to the same connected component, a solution exists
 - 3. Search the solution in the graph
 - 4. Optimize / smooth the path



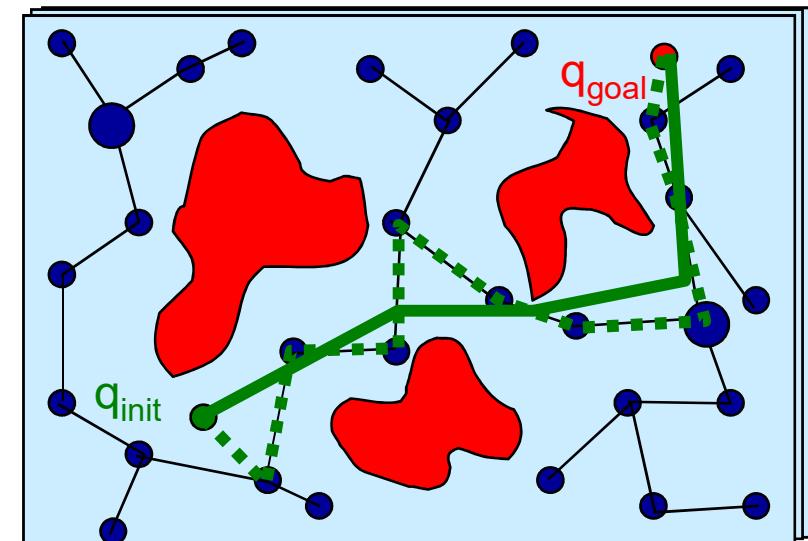
The query phase

- ◆ The PRM is used to solve queries
 - 1. Connect initial and final configurations to the roadmap
 - 2. If nodes belong to the same connected component, a solution exists
 - 3. Search the solution in the graph
 - 4. Optimize / smooth the path



The query phase

- ◆ The PRM is used to solve queries
 - 1. Connect initial and final configurations to the roadmap
 - 2. If nodes belong to the same connected component, a solution exists
 - 3. Search the solution in the graph
 - 4. Optimize / smooth the path

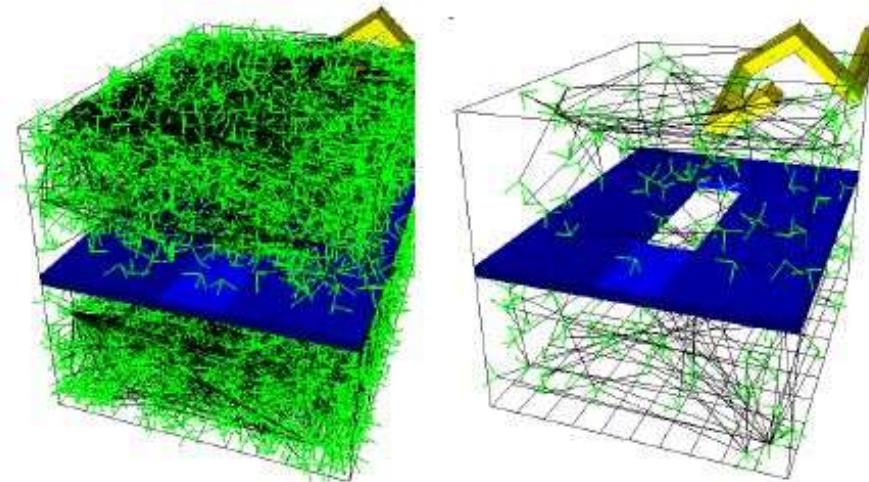


Example



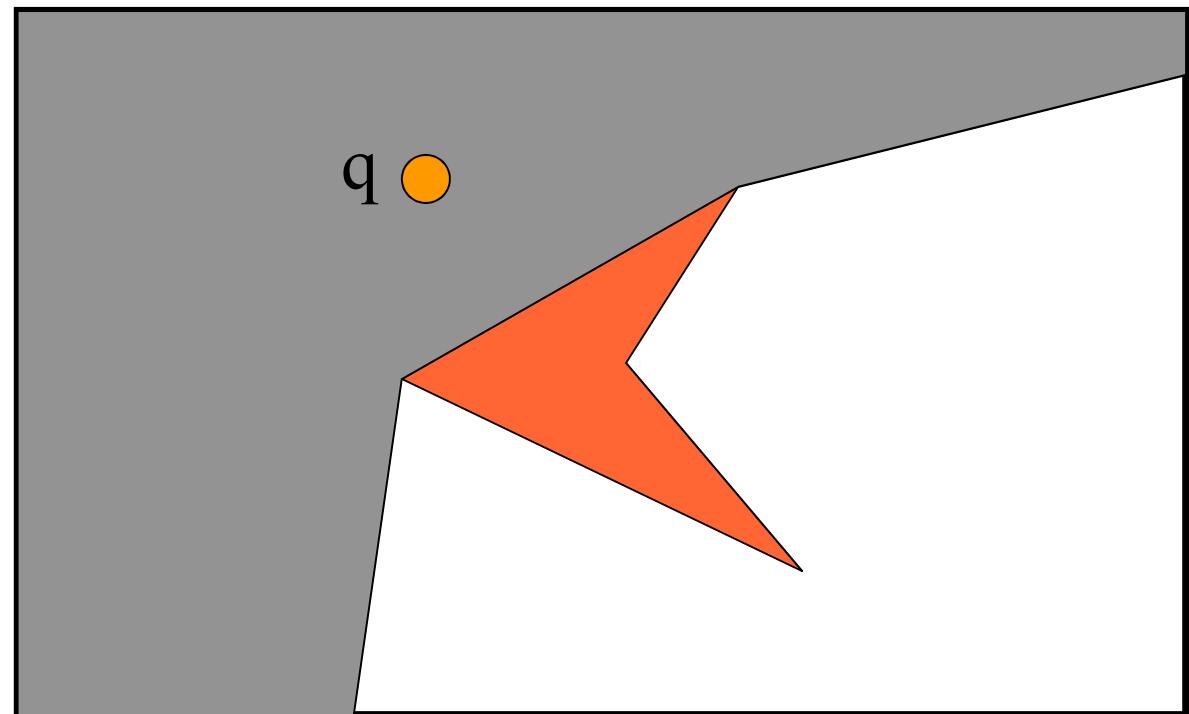
Sampling strategies

- ◆ Planning in very constrained environments generates large roadmaps
 - Construction is time consuming
 - Search is slow
- ◆ Use of sampling strategies
 - Visibility PRMs
 - Gaussian PRMs



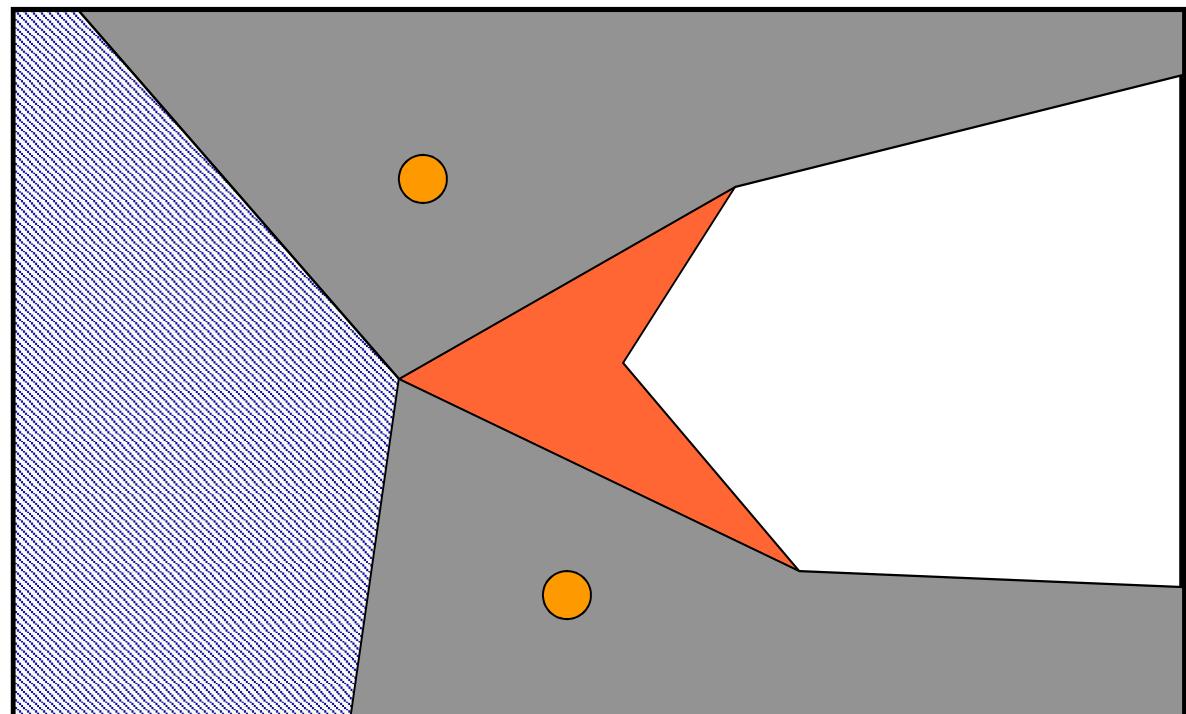
Visibility-PRM

Visibility domain of a configuration q



Visibility-PRM

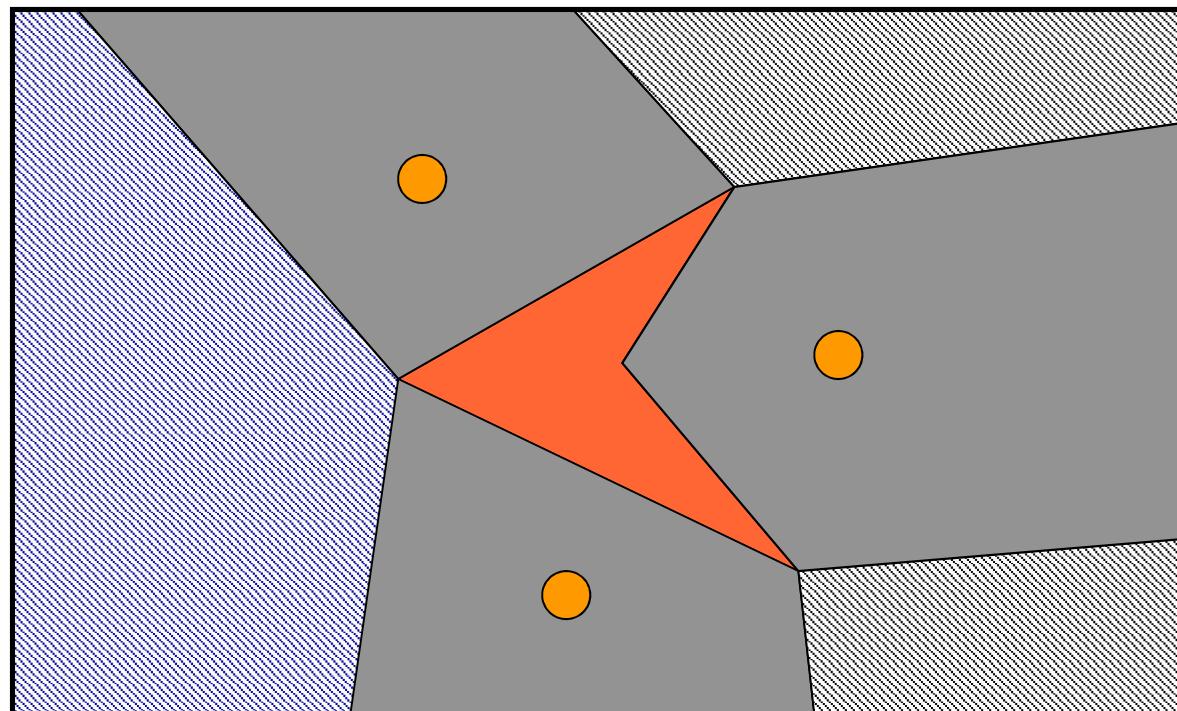
A new configuration is kept if it is outside the visibility domain of other configurations



Visibility-PRM

A new configuration is kept if it is outside the visibility domain of other configurations

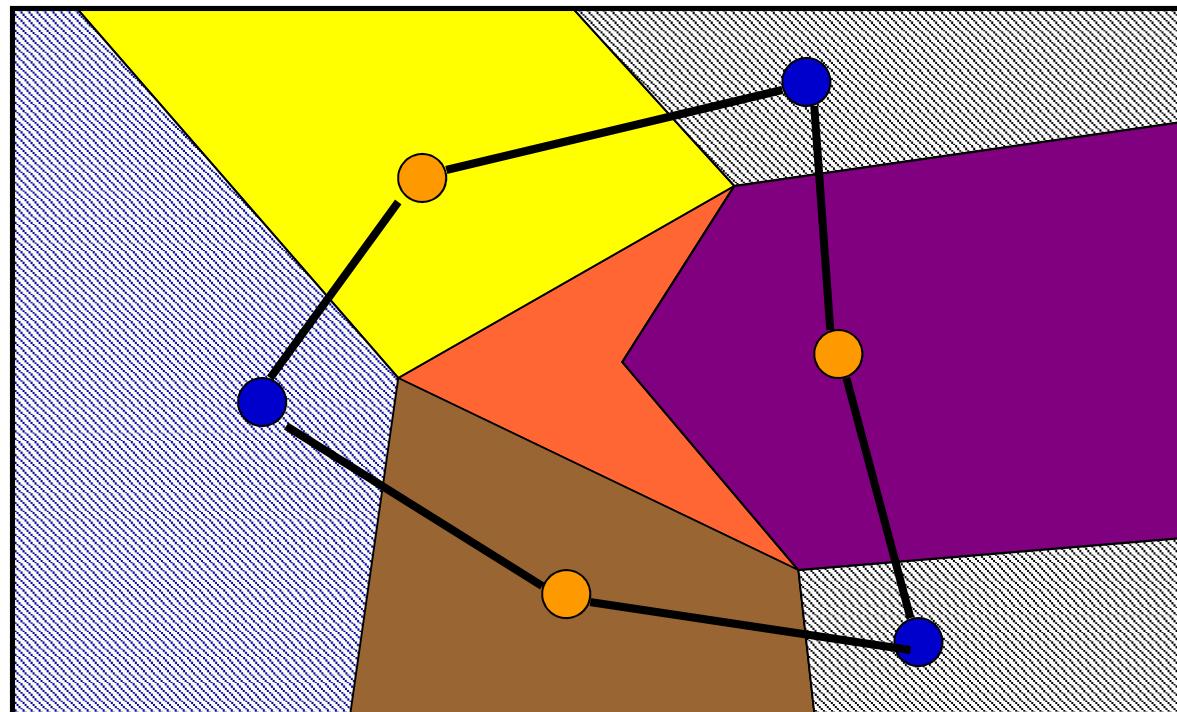
Those configurations are “guardians”



Visibility-PRM

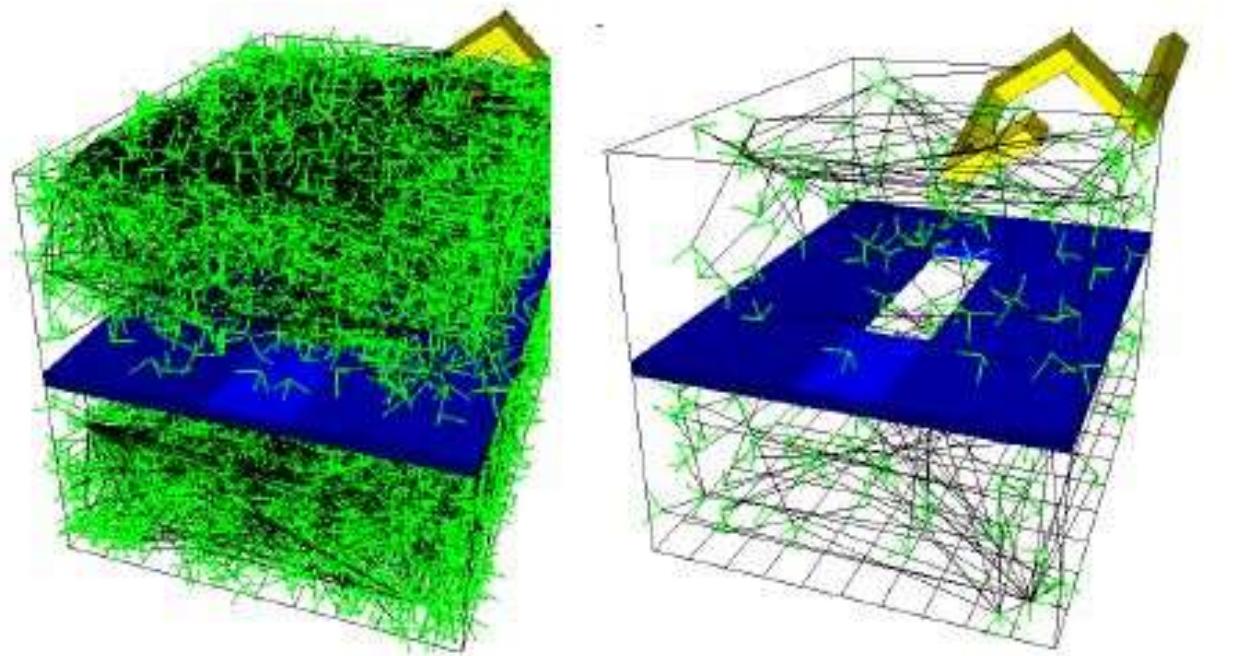
A new configuration is kept if it is outside the visibility domain of other configurations **or if it connects at least two guardians**

Those configurations are “connectors”



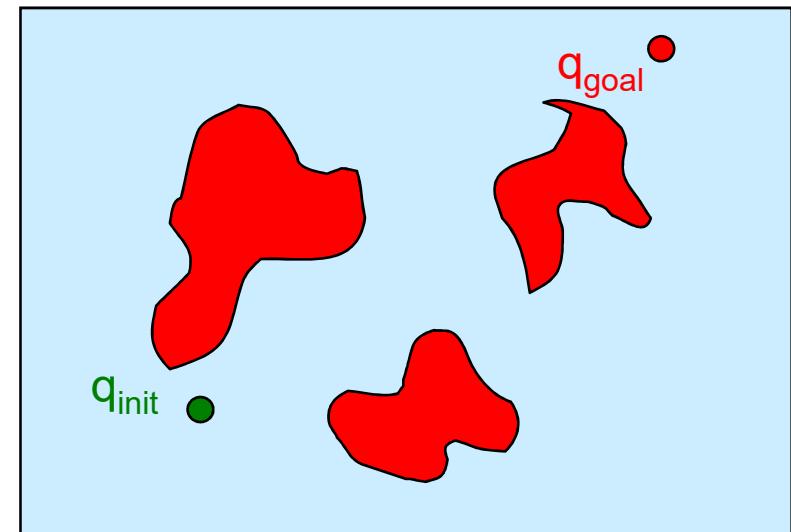
Visibility-PRM

- ◆ Example
 - 6 dimensions configuration space



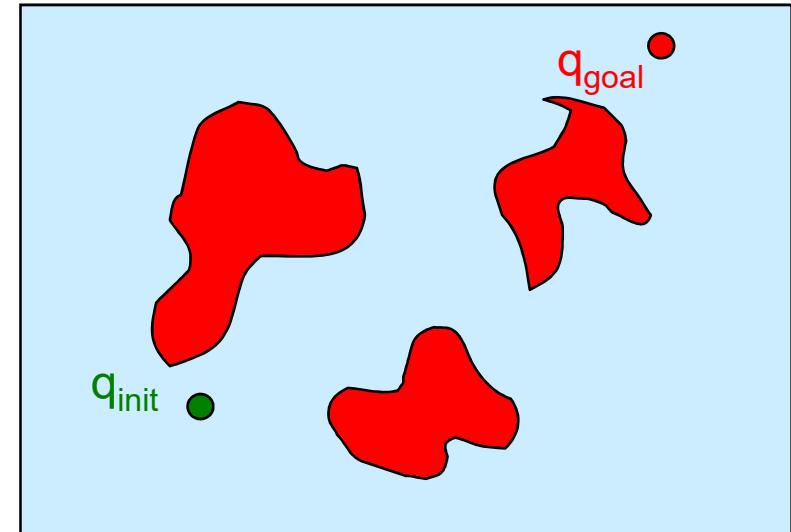
RRT: Rapidly-exploring Random trees

- ◆ PRM is a “multiple queries” method: the same roadmap is used to solve several queries
- ◆ RRT is a “single query” method: a datastructure is computed to solve one query without preliminary exploration of the C-space



RRT: Rapidly-exploring Random trees

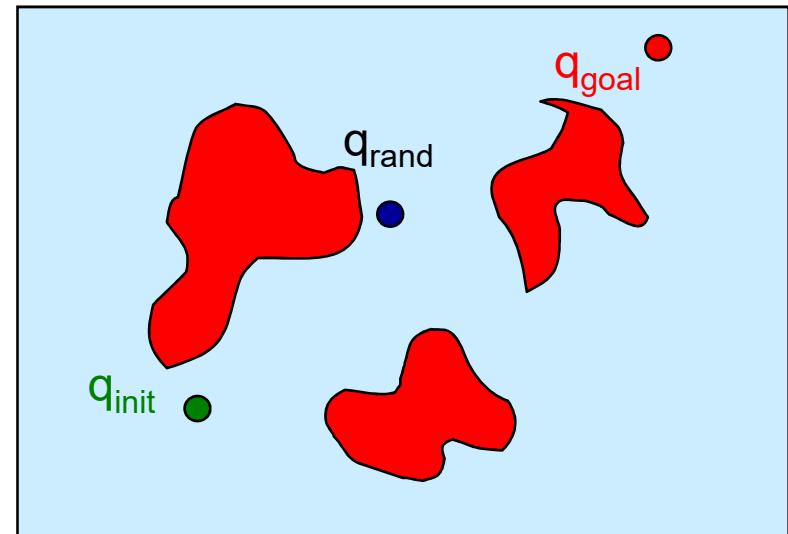
- ◆ An iterative algorithm
 1. Compute q_{rand}
 2. Connect to q_{near}
 3. Insert q_{new}
 4. Go back to 1



RRT: Rapidly-exploring Random trees

- ◆ An iterative algorithm

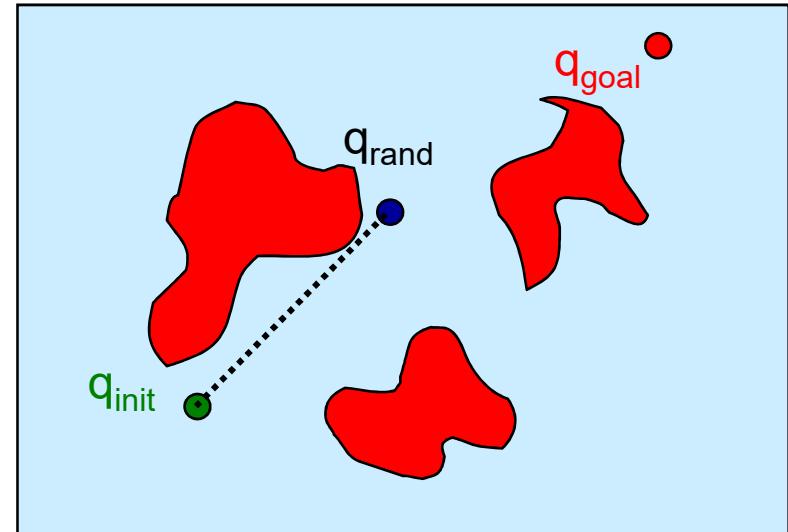
1. Compute q_{rand}
2. Connect to q_{near}
3. Insert q_{new}
4. Go back to 1



RRT: Rapidly-exploring Random trees

- ◆ An iterative algorithm

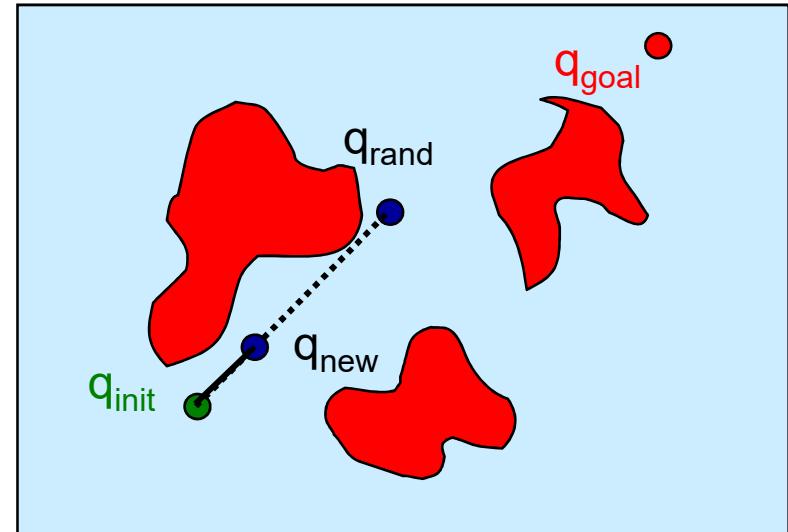
1. Compute q_{rand}
2. Connect to q_{near}
3. Insert q_{new}
4. Go back to 1



RRT: Rapidly-exploring Random trees

- ◆ An iterative algorithm

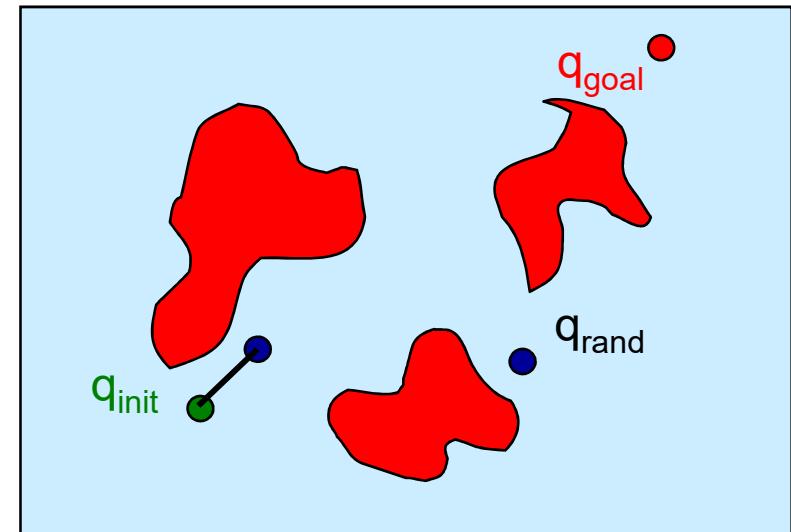
1. Compute q_{rand}
2. Connect to q_{near}
3. Insert q_{new}
4. Go back to 1



RRT: Rapidly-exploring Random trees

- ◆ An iterative algorithm

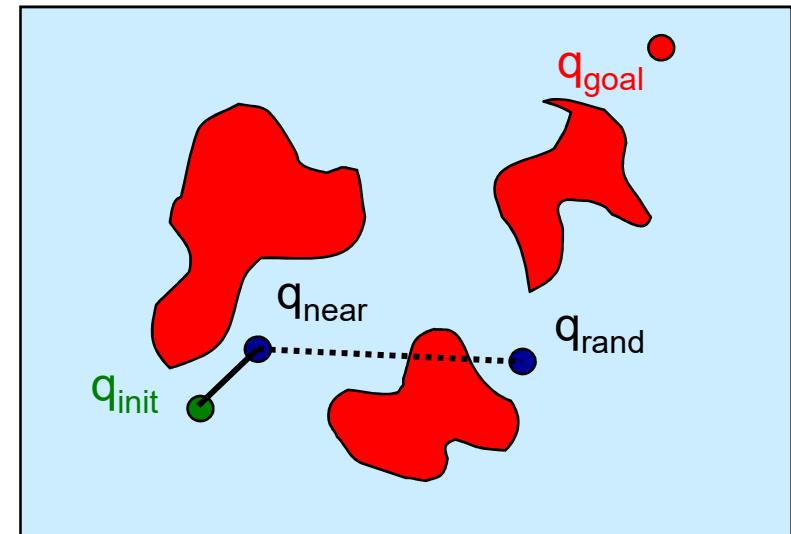
1. Compute q_{rand}
2. Connect to q_{near}
3. Insert q_{new}
4. Go back to 1



RRT: Rapidly-exploring Random trees

- ◆ An iterative algorithm

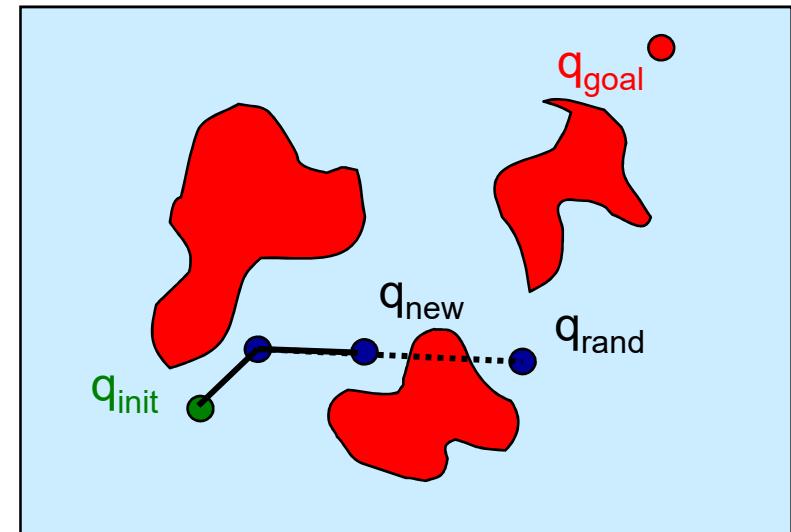
1. Compute q_{rand}
2. Connect to q_{near}
3. Insert q_{new}
4. Go back to 1



RRT: Rapidly-exploring Random trees

- ◆ An iterative algorithm

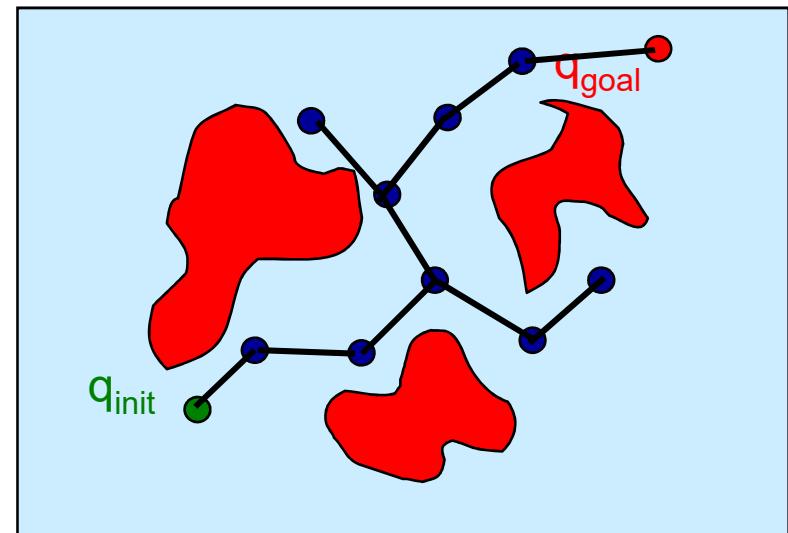
1. Compute q_{rand}
2. Connect to q_{near}
3. Insert q_{new}
4. Go back to 1



RRT:

Rapidly-exploring Random trees

- ◆ An iterative algorithm
 1. Compute q_{rand}
 2. Connect to q_{near}
 3. Insert q_{new}
 4. Go back to 1 until connection between q_{new} and q_{goal}



Exemple

RRT + IK + Motion capture

Synthesizing Animations of Human Manipulation

Katsu Yamane

James Kuffner

Jessica Hodgins

Problem extensions

- ◆ Mobile obstacles
- ◆ Several moving objects
- ◆ Movable objects
- ◆ Deformable objects
- ◆ Partial knowledge of the environment
- ◆ Uncertainty in perception and control
- ◆ Non-holonomic constraints
- ◆ Optimal planning
- ◆ Visibility constraints

Conclusion

- ◆ A large set of techniques have been proposed in motion planning
 - Deterministic search
 - » Approximative representation
 - » Exact representation
 - Potential fields
 - Probabilistic sampling
- ◆ Choice relies on
 - Needed performance
 - Problem dimensions
 - Need of completeness

Conclusion

- ◆ In practice
 - Planning in high dimensional space
 - » Possible in theory
 - » But: memory consumption, computation time...
 - First thing to do: reduce problem dimensionality
 - » Example: pedestrians
 - ◆ Approximation with a cylinder
 - ◆ 2D C-space
 - ◆ Separation between planning / navigation and animation

Search algorithms

Graphs and path planning

- ◆ All planning problems can be represented by a graph
 - Nodes are configurations
 - Edges are paths between configurations
- ◆ *Solving a planning problem is equivalent to finding a sequence of nodes and edges in a graph from one source node to a destination node*
- ◆ *Generally solved by finding the shortest path between a source node and a destination node*
 - *A cost is associated to each edge of the graph*

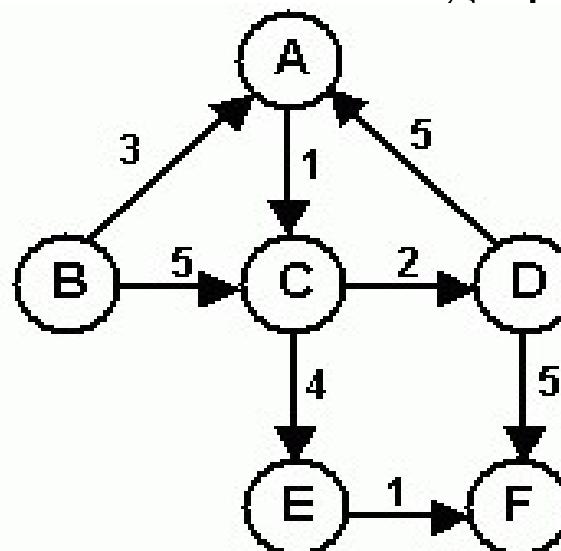
What is the shortest path problem?

- ◆ In an edge-weighted graph, the weight of an edge measures the cost of traveling that edge.
- ◆ For example, in a graph representing a network of airports, the weights could represent: distance, cost or time.
- ◆ Such a graph could be used to answer any of the following:
 - What is the fastest way to get from A to B?
 - Which route from A to B is the least expensive?
 - What is the shortest possible distance from A to B?
- ◆ Each of these questions is an instance of the same problem:

The shortest path problem!

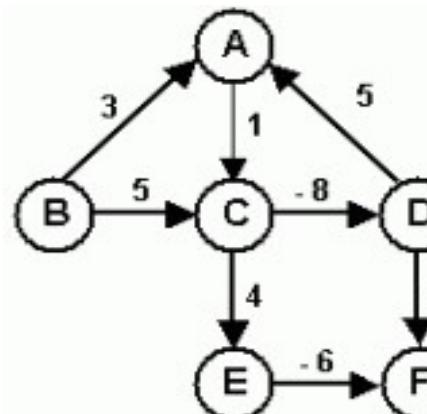
Is the shortest path problem well defined?

- ◆ If all the edges in a graph have non-negative weights, then it is possible to find the shortest path from any two vertices.
- ◆ For example, in the figure below, the shortest path from B to F is { B, A, C, E, F } with a total cost of 9.
- ◆ Thus, the problem is well defined for a graph that contains non-negative weights.



Is the shortest path problem well defined?

- ◆ Things get difficult for a graph with negative weights.
- ◆ For example, the path D, A, C, E, F costs 4 even though the edge (D, A) costs 5 -- the longer the less costly.
- ◆ The problem gets even worse if the graph has a negative cost cycle. e.g. {D, A, C, D}
- ◆ A solution can be found even for negative-weight graphs but not for graphs involving negative cost cycles.



$$\{D, A, C, D, A, C, E, F\} = 2$$

$$\{D, A, C, D, A, C, D, A, C, E, F\} = 0$$

The Dijkstra's algorithm

- ◆ This algorithm solves the single source shortest path problem for a graph with non-negative weights
- ◆ **It finds the shortest path from an initial vertex to all the other vertices**

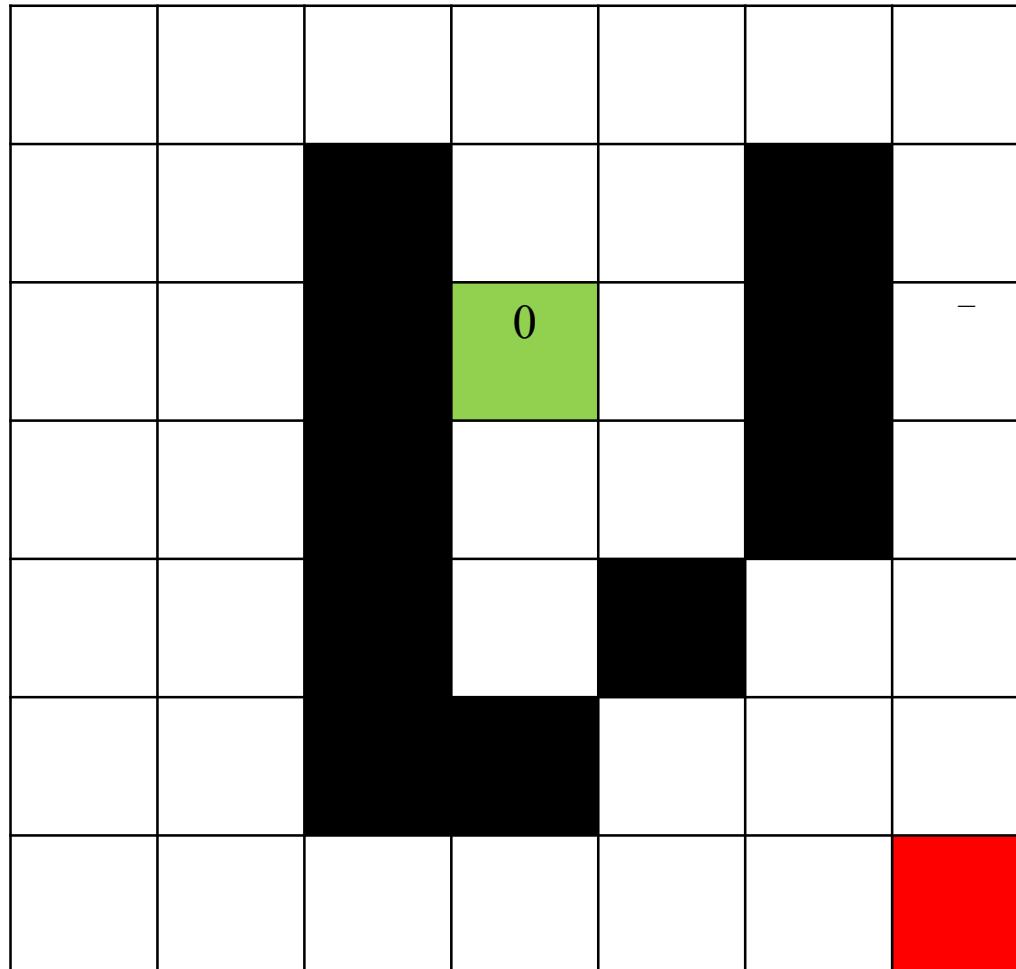
Dijkstra's algorithm: notations

- ◆ **Q** is a priority queue containing pairs (cost, node)
 - This queue is sorted in increasing order of cost
- ◆ **cost(a,b)** is a function returning the cost associated to the edge linking nodes **a** and **b**
- ◆ **cost[a]** is the cost of the optimal path linking **a** to the source node, default value is *infinity*
- ◆ **pre[a]** is the predecessor of **a** on the path to the source node, default value is undefined

Dijkstra's algorithm

```
Q.push(0, source)
while(!Q.empty()) {
    c = Q.pop()
    for all successors S of c {
        if(cost(c,S)+cost[c]<cost[S]) {
            Q.push(cost(c,S)+cost[c], S)
            cost[S] = cost[c] + cost(c,S)
            pre[S]=c
        }
    }
}
```

Example



Example

5	4	3	2	3	4	5
6	5		1	2		6
7	6		0	1		7
8	7		1	2		8
9	8		2		10	9
10	9				11	10
11	10	11				11

The A* algorithm

- ◆ The A* algorithm computes a shortest path from a source vertex to a given destination vertex
 - Faster than Dijkstra as it expands less nodes
- ◆ It uses a heuristic to estimate the cost of reaching the target vertex from the source vertex
- ◆ An **admissible heuristic** is a function that exactly or under estimate the cost of reaching the target from the node
 - *Garanties that the A* algorithm will find the shortest path*

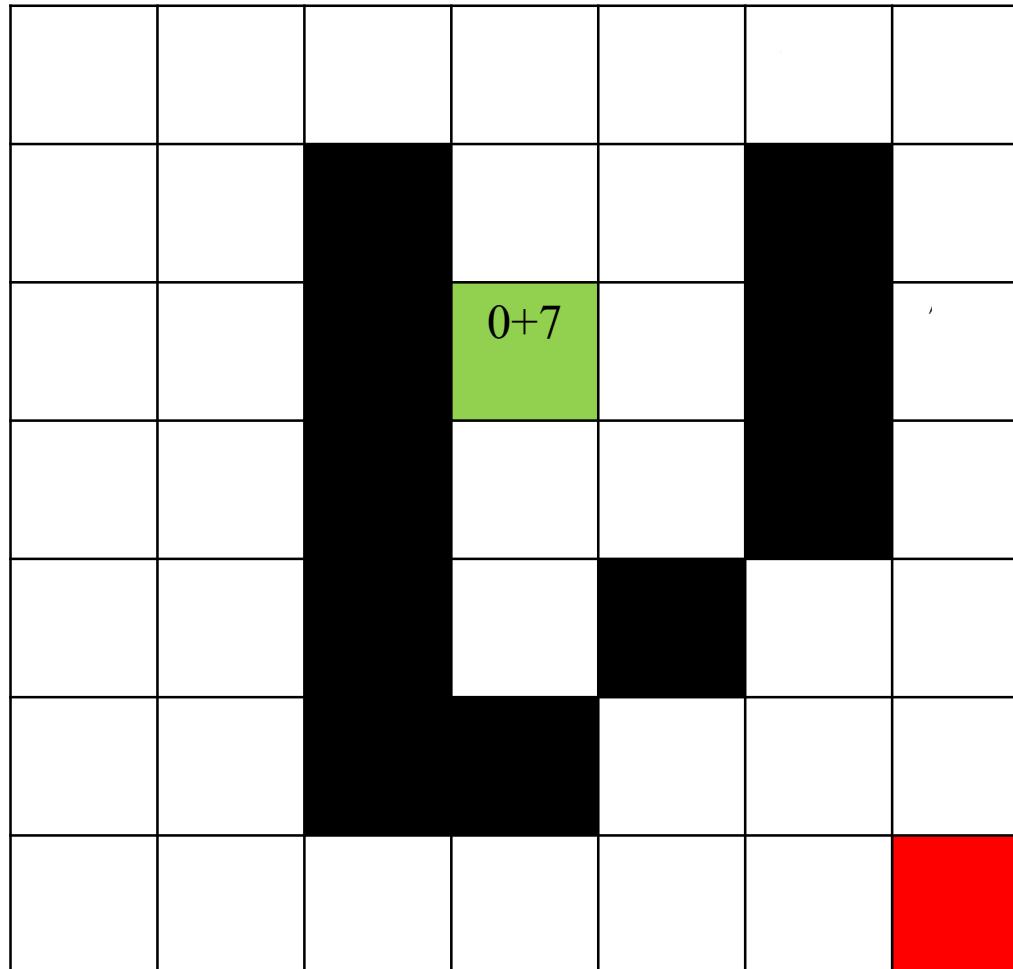
The A* algorithm: notations

- ◆ **Q** is a priority queue containing pairs (cost, node)
 - This queue is sorted in increasing order of cost
- ◆ **cost(a,b)** is a function returning the cost associated to the edge linking nodes a and b
- ◆ **cost[a]** is the cost of the optimal path linking a to the source node, default value is *infinity*
- ◆ **pre[a]** is the predecessor of a on the path to the source node, default value is undefined
- ◆ **H(a)** is the heuristic function evaluating the cost from vertex a to the destination vertex

The A* algorithm

```
Q.push(H(source), source)
while(!Q.empty()) {
    c = Q.pop()
    if(c==target) { /* Path is found ! */ }
    for all successors S of c {
        if (cost(c,S)+cost[c]<cost[S]) {
            Q.push(cost(c,S)+cost[c]+H(S), S)
            cost[S] = cost[c] + cost(c,S)
            pre[S]=c
        }
    }
}
```

Example



Example

			$2+9$	$3+8$	$4+7$	$5+6$
			$1+8$	$2+7$		$6+5$
			$0+7$	$1+6$		$7+4$
			$1+6$	$2+5$		$8+3$
			$2+5$		$10+3$	$9+2$
					$11+2$	$10+1$
						$11+0$

Comparison between Dijkstra and A*

5	4	3	2	3	4	5				2+9	3+8	4+7	5+6
6	5		1	2		6				1+8	2+7		6+5
7	6		0	1		7				0+7	1+6		7+4
8	7		1	2		8				1+6	2+5		8+3
9	8		2		10	9				2+5		10+3	9+2
10	9				11	10						11+2	10+1
11	10	11				11							11+0

The weighted A*

- ◆ Let define $H'(n) = \epsilon H(N)$ with $\epsilon > 1$
- ◆ If $H'(n)$ is used as a heuristic in the A*
 - Produced path is suboptimal and is no more than ϵ times longer than the optimal path
 - Can speed up the algorithm by exploring less nodes

Impact of the heuristic function

- ◆ If $H(n)=0$, A* turns into Dijkstra algorithm
- ◆ If $H(n)$ is always lower or equal to the cost of moving from n to the target, A* is guaranteed to find the shortest path
- ◆ If $H(n)$ is always equal to the cost of moving from n to the target, A* exactly follows the shortest path
- ◆ If $H(n)$ is sometimes greater than the cost of moving from n to the target, A* is not guaranteed to find a shortest path but can run faster

The IDA* algorithm

- ◆ The IDA* algorithm computes a shortest path from a source vertex to a given destination vertex
- ◆ IDA* is a variant of the A* search which uses iterative deepening to keep memory usage lower than A*
- ◆ Suitable for the exploration of very large graphs that do not fit into memory
 - Ex : implicit graphs generated by the exploration of possible states

The IDA* algorithm

```
IDA_Star(source, target) {
    cost_limit = H(source)
    solution = None
    while(true) {
        (solution, cost_limit) = DFS(0, source, cost_limit, [source])
        if(solution!=None) { return (solution, cost_limit) } /* Solution found */
        if(cost_limit==infinity) { return None } /* No solution exist */
    }
}
```

The IDA* algorithm

```
DFS(start_cost, node, cost_limit, path) {  
    mini_cost = start_cost + H(node)  
    if(mini_cost>cost_limit) { return (None, mini_cost) }  
    if(node==target) { return (path, cost_limit) }  
    next_cost_limit = infinity  
    for all successors S of node {  
        newStartCost = start_cost+cost(node, S)  
        (solution, new_cost_limit) = DFS(newStartCost, S, cost_limit, path+[S])  
        if(solution!=None) { return (solution, new_cost_limit) }  
        next_cost_limit = min(next_cost_limit, new_cost_limit)  
    }  
    return (None, next_cost_limit)  
}
```

Example

- ◆ Pass 1: cost_limit = 7
 - new limit = 9

A 8x8 grid representing a search space. The grid contains several colored cells indicating different states or values:

- Black Cells:** These represent obstacles or不可达 states. They form vertical and horizontal barriers.
- Red Text:** Values labeled in red are 1+8 and 2+7, located in the top-right corner.
- Green Cell:** A single cell at row 4, column 5 is highlighted green, containing the value 0+7.
- White Cells:** Most cells are white, representing open or unexplored states.
- Red Cell:** A single cell at row 7, column 8 is filled red.
- Other Values:** Some cells contain values like 1+6, 2+5, and 2+6, which likely represent f-values or g-values from a search algorithm like A*.

Example

- ◆ Pass 1: cost_limit = 7
 - new limit = 9
- ◆ Pass 2 : cost_limit = 9
 - new limit = 11

				2+9	3+8		
				1+8	2+7		
				0+7	1+6		
				1+6	2+5		
				2+5			
							90

Example

- ◆ Pass 1: cost_limit = 7
 - new limit = 9
- ◆ Pass 2 : cost_limit = 9
 - new limit = 11
- ◆ Pass 3: cost_limit = 11
 - Goal reached

		3+ 10	2+9	3+8	4+7	5+6
			1+8	2+7		6+5
			0+7	1+6		7+4
			1+6	2+5		8+3
			2+5		10+ 3	9+2
					11+ 2	10+ 1
						11+ 0

IDA* algorithm

- ◆ Without enhancements
 - Slower than A*
 - » A* uses a sorted Q => impact on performances
 - » IDA* explores several times the same configurations with different paths

- ◆ Possible enhancements
 - Do not explore nodes already present in the current path
 - Using a cache to store best estimated path lengths in some configurations

Heuristics in path planning

- ◆ In grids

- Manhattan distance

- » Useful in case of 4-connexity

- » $h(n) = D * (\text{abs}(n.x - \text{goal}.x) + \text{abs}(n.y - \text{goal}.y))$

- » D: min cost of a move to an adjacent cell

- Diagonal distance V1

- » Useful in case of 8-connexity

- » $h(n) = D * \max(\text{abs}(n.x - \text{goal}.x), \text{abs}(n.y - \text{goal}.y))$

- » D: min cost of a move to an adjacent cell

Heuristics in path planning

- ◆ In grids

- Diagonal distance V2

- » $h_{\text{diagonal}}(n) = \min(\text{abs}(n.x - \text{goal}.x), \text{abs}(n.y - \text{goal}.y))$

- » $h_{\text{straight}}(n) = (\text{abs}(n.x - \text{goal}.x) + \text{abs}(n.y - \text{goal}.y))$

- » $h(n) = D2 * h_{\text{diagonal}}(n) + D * (h_{\text{straight}}(n) - 2 * h_{\text{diagonal}}(n))$

- » $D2$: cost of a diagonal move

- » D : cost of a vertical / horizontal move

- ◆ In continuous space / grids

- Euclidian distance

- » $h(n) = D * \sqrt{((n.x - \text{goal}.x)^2 + (n.y - \text{goal}.y)^2)}$

Which algorithm for what purpose?

- ◆ Dijkstra
 - Useful to compute all shortest paths to a given target from all nodes
- ◆ A*
 - Useful to compute the shortest path from a source to a destination
- ◆ IDA*
 - Useful to compute the shortest path from a given source to a given target if graph is implicitly generated and cannot be stored explicitly

Discussions on examples

- ◆ Which algorithm and (eventually) which heuristic?
- ◆ A set of robots should reach a unique target
 - Environment is represented with a grid
- ◆ Flee from a set of dangerous positions
 - Environment is represented with a roadmap
- ◆ Find the nearest cash dispenser knowing several cash dispensers positions in a 2D world
 - Environment is represented with a roadmap

Other kind of algorithms

- ◆ Incremental replanning algorithms
 - Handle changes in costs associated to the graph
 - Ex: D*, focussed D* and D* lite algorithms
- ◆ Anytime algorithms
 - Agent must react quickly and the planning problem is complex
 - Find an approximate solution, then refine it
 - ARA* (Anytime repairing A*)
- ◆ Dave Ferguson, Maxim Likhachev, and Anthony Stentz. A Guide to Heuristic-based Path Planning. ICAPS 2005.

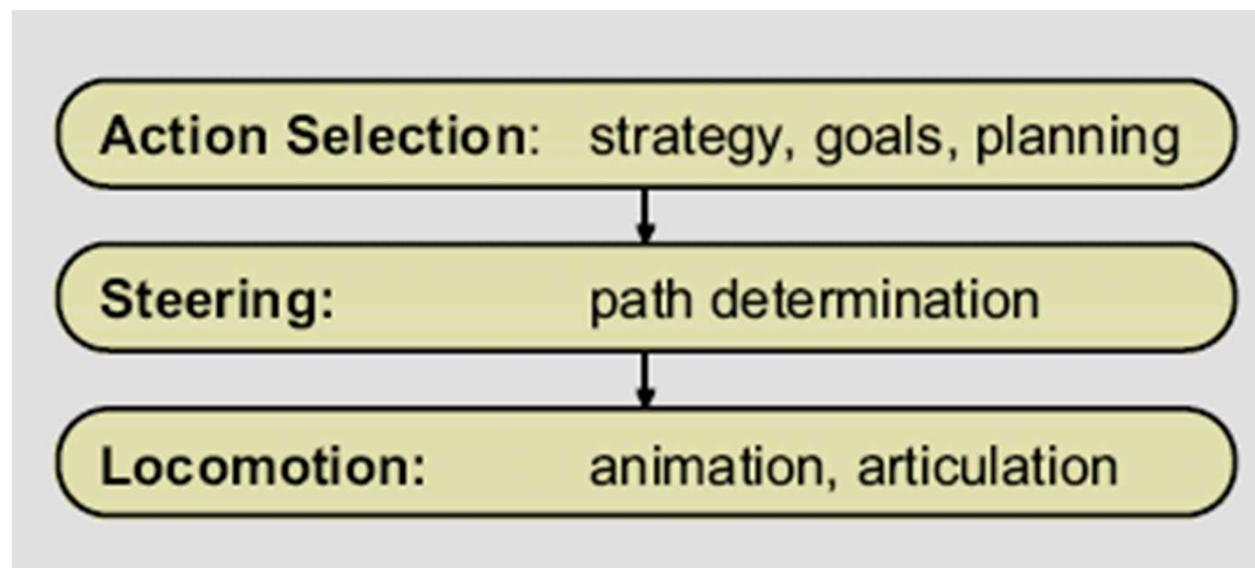
http://www.ri.cmu.edu/publication_view.html?pub_id=5085

Steering behaviors and group movements

Introduction

- ◆ Autonomous Characters
 - Autonomous agents
 - NPCs in Games (Non-player characters)
- ◆ Applications
 - Robotics
 - Artificial Intelligence (AI)
 - Artificial Life
- ◆ References
 - Craig W. Reynolds
 - 1987 “Flocks, Herds, and Schools: A Distributed Behavioral Model”, Siggraph’87 Proceedings
 - 1999 “Steering Behaviors for Autonomous Characters”, GDC Proceedings*
 - www.red3d.com/cwr/steer

Motion behavior



Steering

- ◆ Path Determination
 - Path finding or path planning
- ◆ Behaviors
 - Seek & flee
 - Pursuit & evasion
 - Obstacle avoidance
 - Wander
 - Path following
 - Unaligned collision avoidance
- ◆ Group steering
 - Separation, cohesion, alignment

A simple vehicle model

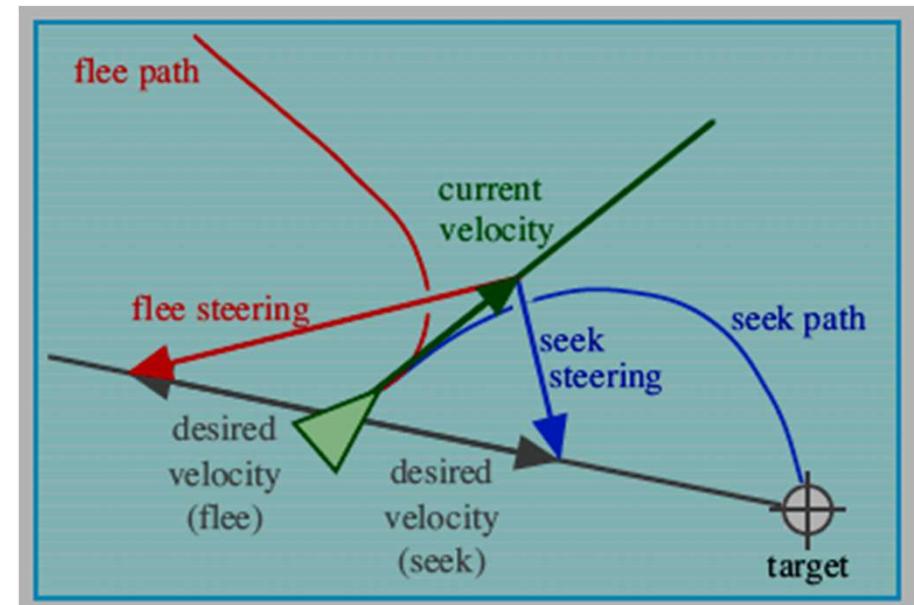
- ◆ A Point Mass
 - Linear momentum
 - No rotational momentum
- ◆ Parameters
 - Mass (scalar)
 - Position (vector)
 - Velocity (vector)
 - » Modified by applied forces
 - Max speed (scalar)
 - » Top speed of a vehicle
 - Max steering force (scalar)
 - » Self-applied
 - Orientation (scalar)
 - » Car / aircraft

Computation of a new position: Euler integration

- ◆ $\text{steer_force} = \text{Truncate}(\text{steering}, \text{Max_force})$
 - ◆ $\text{Acceleration} = \text{steer_force} / \text{mass}$
 - ◆ $\text{Velocity} = \text{Truncate}(\text{Old_velocity} + \text{Acceleration} * \text{Time_step}, \text{Max_speed})$
 - ◆ $\text{New_position} = \text{Old_position} + \text{Velocity} * \text{Time_step}$
-
- Steering: difference between current velocity and desired velocity
 - Truncate: a function that limits the norm of the vector to a maximum value
 - Time_step: time elapsed since last update

Seek and flee behaviors

- ◆ Seek: pursuit to a **Static Target**
 - Steer a character toward to a target position
- ◆ Seek Steering force
 - $\text{desired_velocity} = \text{normalize}(\text{target} - \text{position}) * \text{max_speed}$
 - $\text{steering} = \text{desired_velocity} - \text{velocity}$
- ◆ Flee
 - Inverse of Seek
- ◆ Variants
 - Arrival
 - Pursuit of a moving target

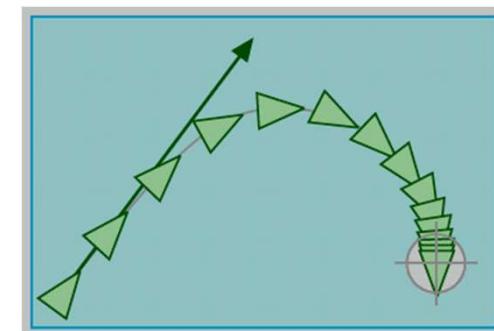


Arrival behavior

◆ A Stopping Radius

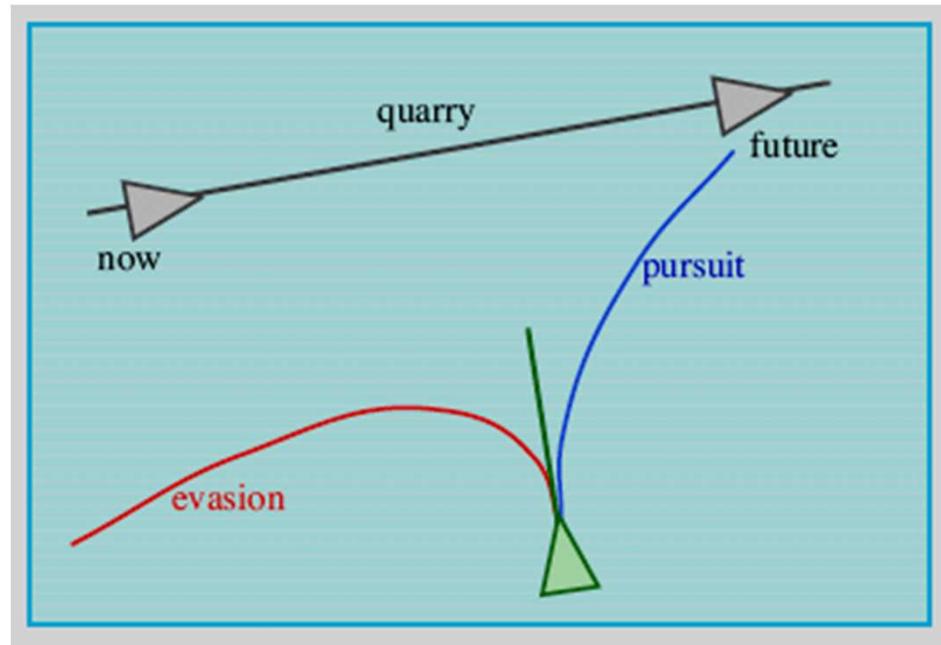
- Outside the radius, arrival is identical to seek
- Inside the radius, the speed is decreased down to zero
 - » $\text{target_offset} = \text{target} - \text{position}$
 - » $\text{distance} = \text{length}(\text{target_offset})$
 - » $\text{ramped_speed} = \text{max_speed} * (\text{distance}/\text{slowing_distance})$
 - » $\text{clipped_speed} = \min(\text{ramped_speed}, \text{max_speed})$
 - » $\text{desired_velocity} = (\text{clipped_speed}/\text{distance}) * \text{target_offset}$
 - » $\text{steering} = \text{desired_velocity} - \text{Velocity}$

➤ *Rq: slowing_distance is a parameter of the behavior*



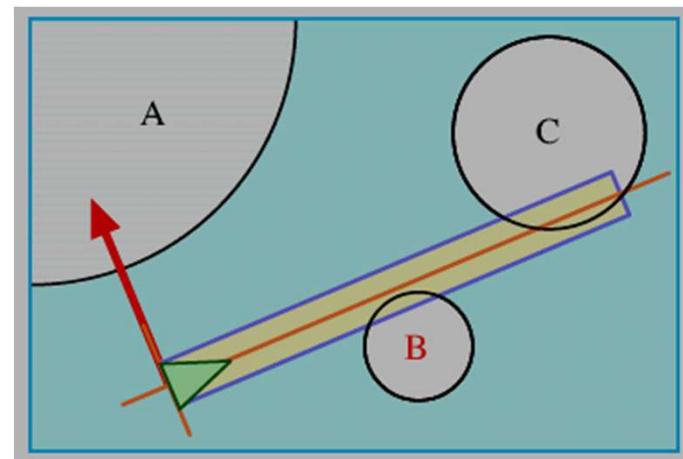
Pursuit and evasion behaviors

- ◆ Another character is moving
- ◆ Pursuit: apply seek behavior to a character's predicted position
- ◆ Evasion: apply flee behavior to a character's predicted position
- ◆ Position prediction
 - Ex: linear extrapolation based of oberved velocity vector and a given extrapolation time



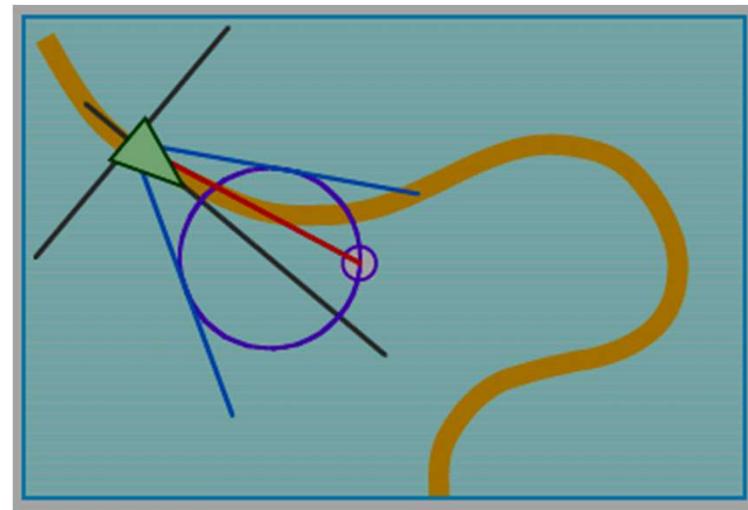
Obstacle avoidance behavior

- ◆ Use a bounding sphere / circle for the controlled character
- ◆ Obstacle selector
 - All points which distance to a segment is lower than a given threshold
 - Segment: (position, position+velocity*time)
 - Distance threshold: bounding sphere / circle radius



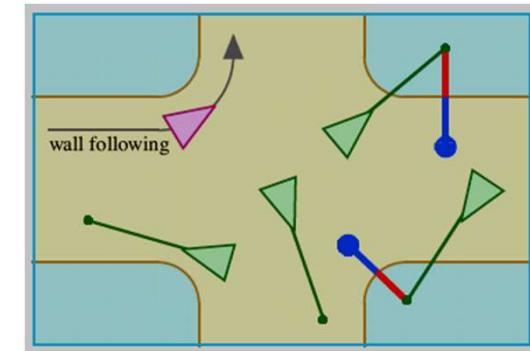
Wander behavior

- ◆ Random steering
- ◆ One Solution :
 - Retain steering direction state
 - » Constrain steering force to the sphere surface located slightly ahead of the character
 - Make small random displacements of the steering at each frame
 - » A small sphere on sphere surface to indicate and constrain the displacement

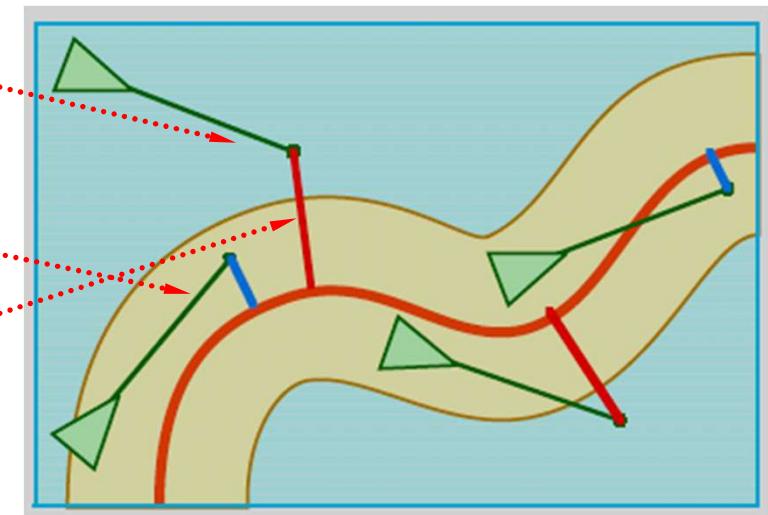


Path following behavior

- ◆ A Path
 - Spine
 - » A spline or poly-line to define the path
 - Pipe
 - » The tube or generated cylinder by a defined “radius”

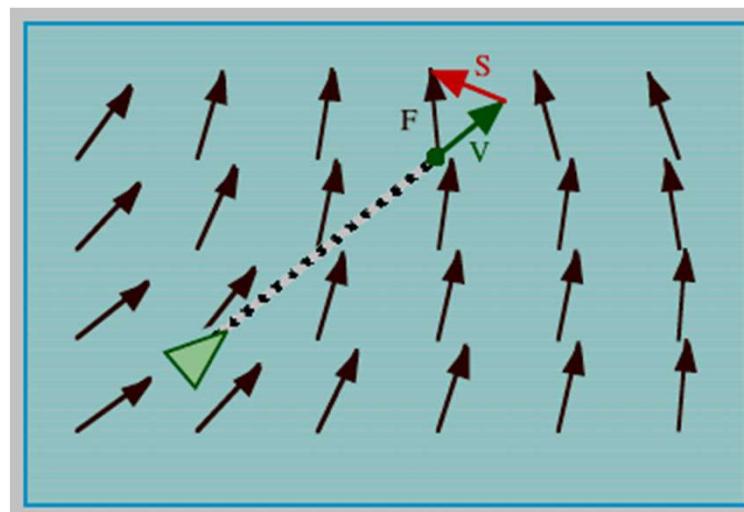


- ◆ Following
 - A velocity-based prediction position
 - » Inside the tube
 - ◆ Do nothing about steering
 - » Outside the tube
 - ◆ “Seek” to the on-path projection



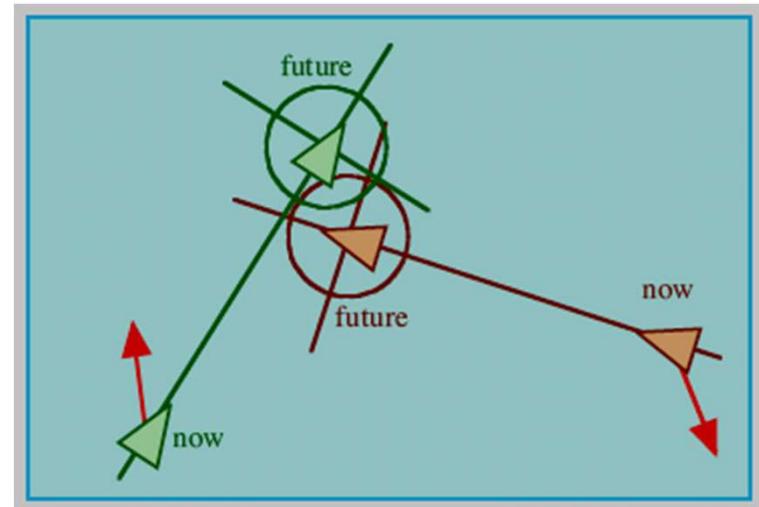
Flow field following behavior

- ◆ A Flow Field Environment is Defined
- ◆ Computation of the steering force
 - The next position of the character is estimated
 - The flow field is sampled at this position (flow_velocity)
 - Steering = flow_velocity - velocity



Unaligned collision avoidance behavior

- ◆ Turn Away from Possible Collision
- ◆ Predict the Potential Collision
 - Use bounding spheres
- ◆ If possibly collide,
 - Apply the steering on both characters
 - Steering direction is possible collision result
 - » Use “future” possible position
 - » The connected line between the two sphere centers



Steering of groups of characters

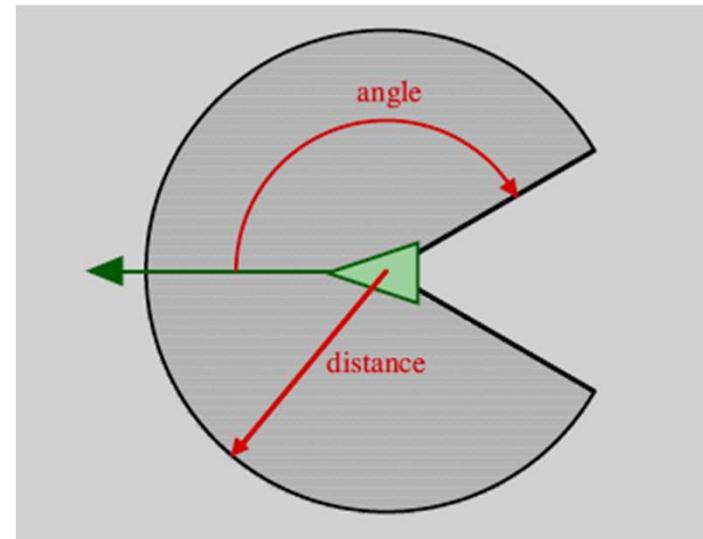
- ◆ Steering behaviors determining how the character reacts to the other characters within his local neighborhood

- ◆ Behaviors include
 - Separation
 - Cohesion
 - Alignment

Local neighbourhood

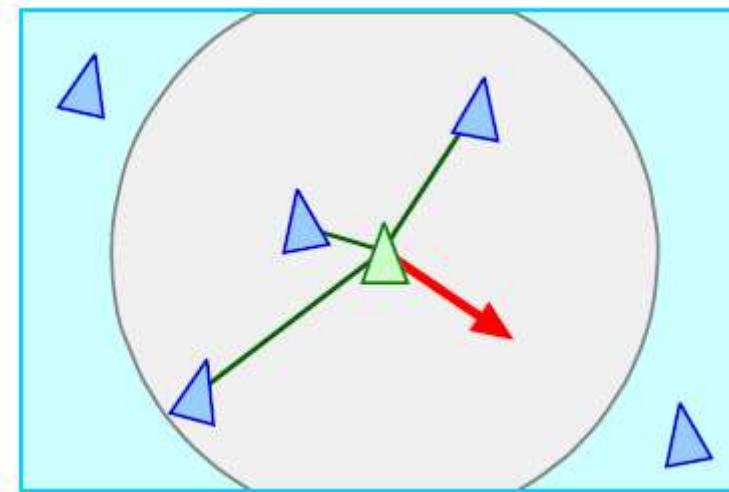
- ◆ The local neighborhood is defined with
 - A distance
 - A field-of-view (angle)

- ◆ Need of a spatial data structure for neighborhood retrieval
 - Kd-Tree
 - Octree / Quadtree
 - Regular grids
 - ...



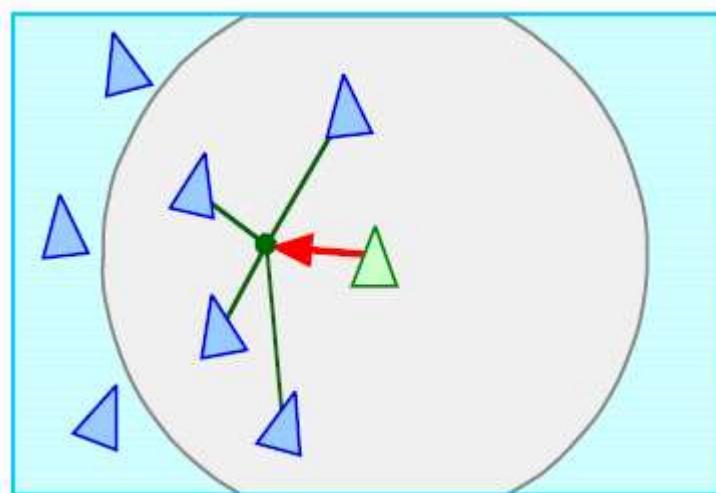
Separation behavior

- ◆ The character maintains a distance from others
 - Compute the repulsive forces within local neighborhood
 - » Calculate the position vector for each nearby
 - » Normalize it
 - » Weight the magnitude with distance
 - ◆ $1/\text{distance}$
 - » Sum the result forces
 - » Negate it



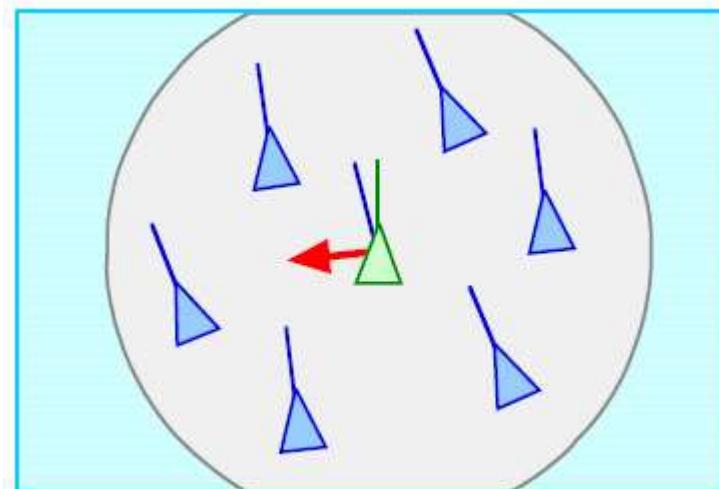
Cohesion behavior

- ◆ The character navigates among the others
 - Compute the cohesive forces within local neighborhood
 - » Compute the average position of the other nearbys
 - ◆ Gravity center
 - » Apply “Seek” to the position



Alignment behavior

- ◆ The character to aligns with the others
 - Compute the steering force
 - » Average the velocity of all other characters nearby
 - » The result is the desired velocity
 - » Correct the current velocity to the desired one with the steering force



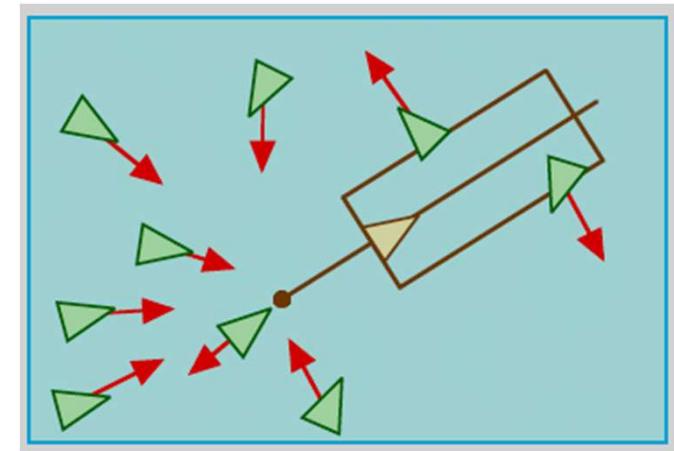
Flocking behaviors

- ◆ Boids Model of Flocks
 - [Reynolds 87]
- ◆ Combination of
 - Separation steering
 - Cohesion steering
 - Alignment steering
- ◆ For Each Combination
 - A weight for combining
 - A distance
 - An Angle

Leader following behavior

- ◆ Follow a Leader

- Stay with the leader
 - » “Pursuit” behavior (Arrival style)
- Stay out of the leader’s way
 - » Defined as “next position” with an extension
 - » “Evasion” behavior when inside the above area
- “Separation” behavior for the followers



Steering behavior

- ◆ A Simple Vehicle Model with Local Neighborhood
- ◆ Common Steering Behaviors
 - Seek, Flee, Pursuit, Evasion
 - Offset pursuit, Arrival, Obstacle avoidance
 - Wander, Path following, Wall following
 - Flow field following, Unaligned collision avoidance
 - Separation, Cohesion, Alignment
 - Flocking, Leader following
- ◆ Combination of the behaviors defined above
 - Intelligent steering behavior

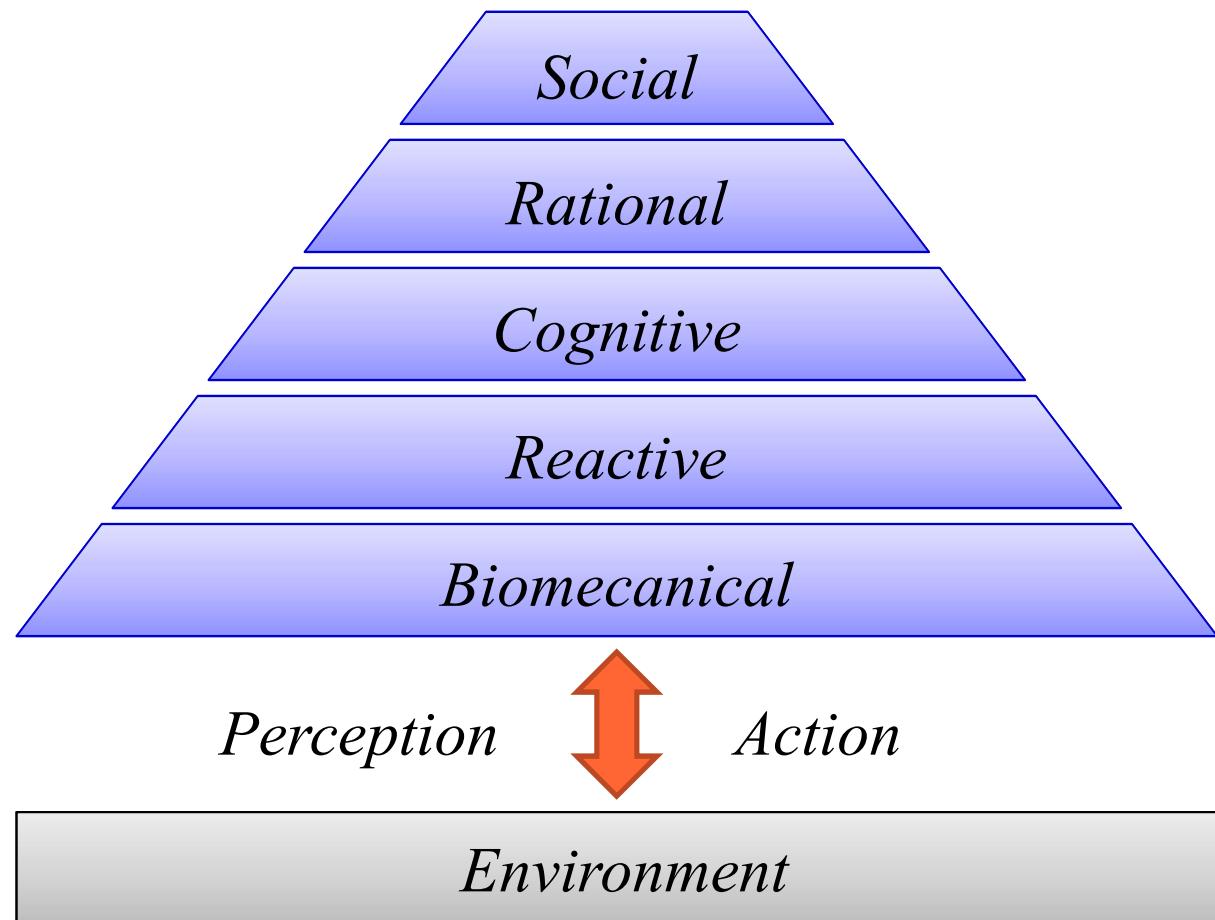
Behavioral animation

Behavior modeling

Behavioral animation

- ◆ Goal:
 - Populating virtual environments
 - Interactive applications
 - » Interactions not always predetermined
 - » Automatic adaptation of the characters' behavior
 - Simulation
 - » Crowd simulation: environment validation, security
 - Automating creation processes
 - » Crowds in films
 - » Golaem, RTK Crowds, MASSIVE

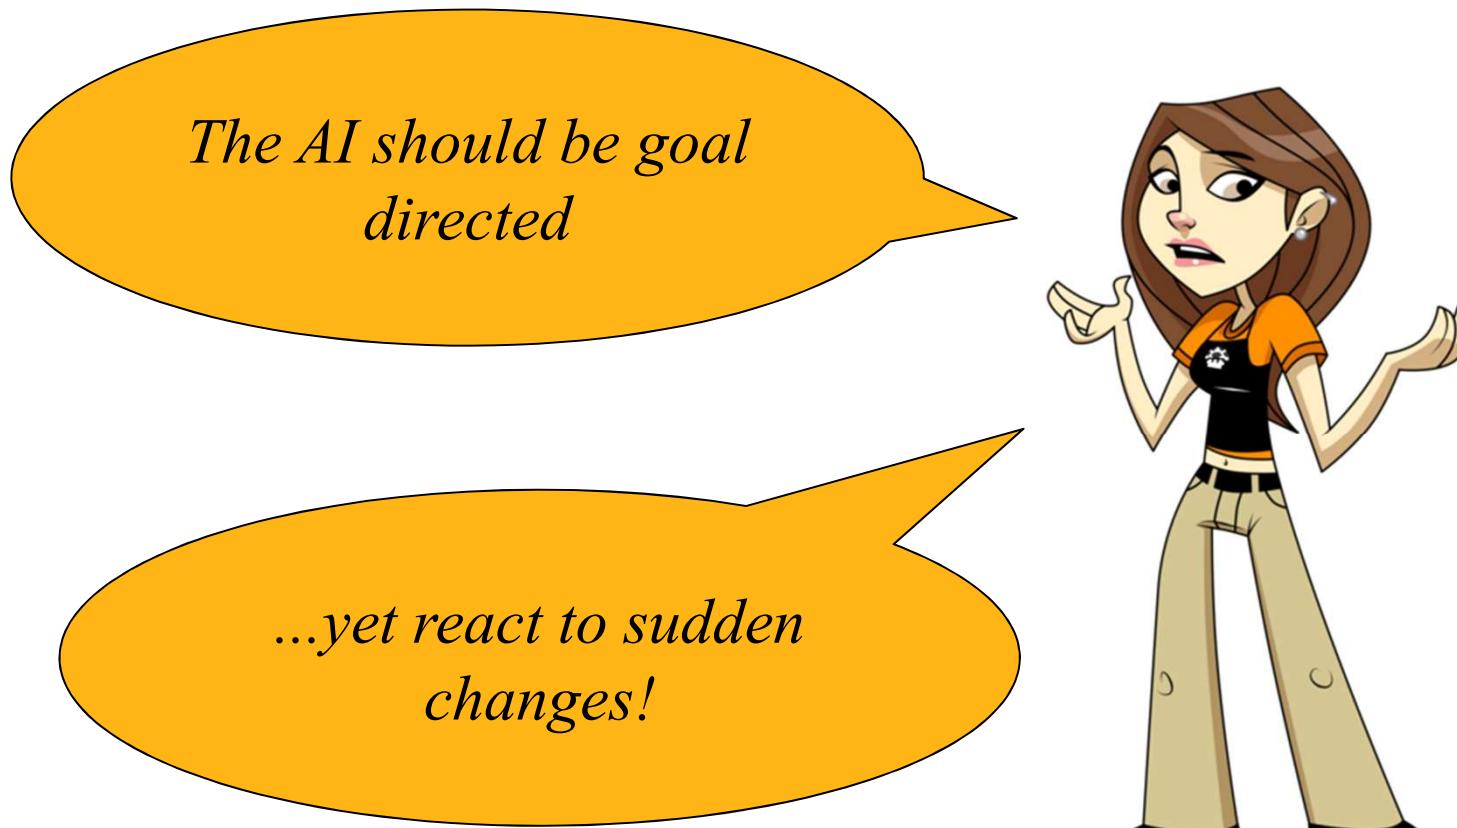
Behavior pyramid



Wish list...



Wish list...



Challenges

- ◆ Big environments
- ◆ Large number of entities
- ◆ Advanced animation & control
- ◆ Designer supervision
 - But AI behave autonomously
- ◆ AI should be goal directed
 - But react to sudden changes

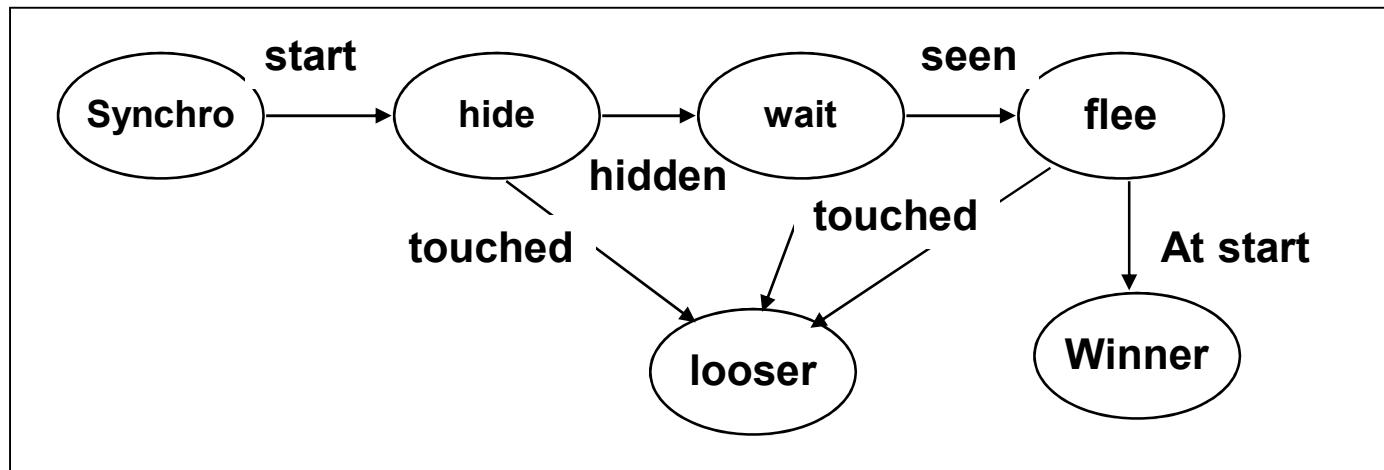


Rule based systems

- ◆ General form:
 - If **condition** then **action**
 - Ex :
 - » If near enemy and healthy then attack
 - » If near enemy and not healthy then flee
 - Actions can be atomic or fire a task to complete...
- ◆ advantage: simple approach, easy to describe
- ◆ Problems:
 - » Conflicts between rules
 - ◆ Need of a strategy, in general priorities
 - » Do not ensure temporal consistency of behavior

Finite state machines

- ◆ Example: hide-and-seek game



Finite state machines

- ◆ Composed of
 - States
 - » An atomic activity
 - » Entry / exit state actions
 - » While in state, activity
 - Transitions
 - » Conditional change of activity
 - » A condition
 - » Possible action on transition
- ◆ State machine = logical suit of activities
- ◆ One initial state, Several ending state

Finite states machines

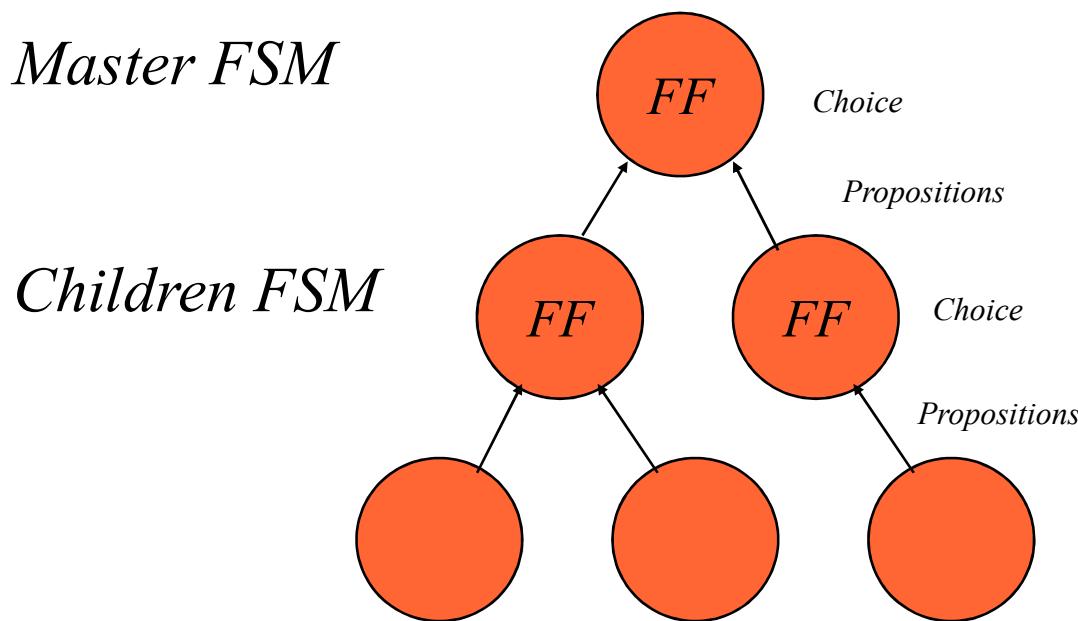
- ◆ State machine stacks
 - One state machine can call another
 - Current state machine is pushed, the other is ran
 - Once the other finishes, the calling state machine is pop and continue its execution
 - Transposition of the notion of procedure call for state machines
 - » A state in the upper level is described through a state machine at the lower level

Finite state machines

- ◆ Parallel state machines
 - Several state machines can run in parallel
 - Use semaphores and message passing for synchronization and communication
 - » Warning: dead locks...
 - Advantage: simulate the product of the two state machines
 - Example: one state machine for displacement, another to shoot

Finite state machines

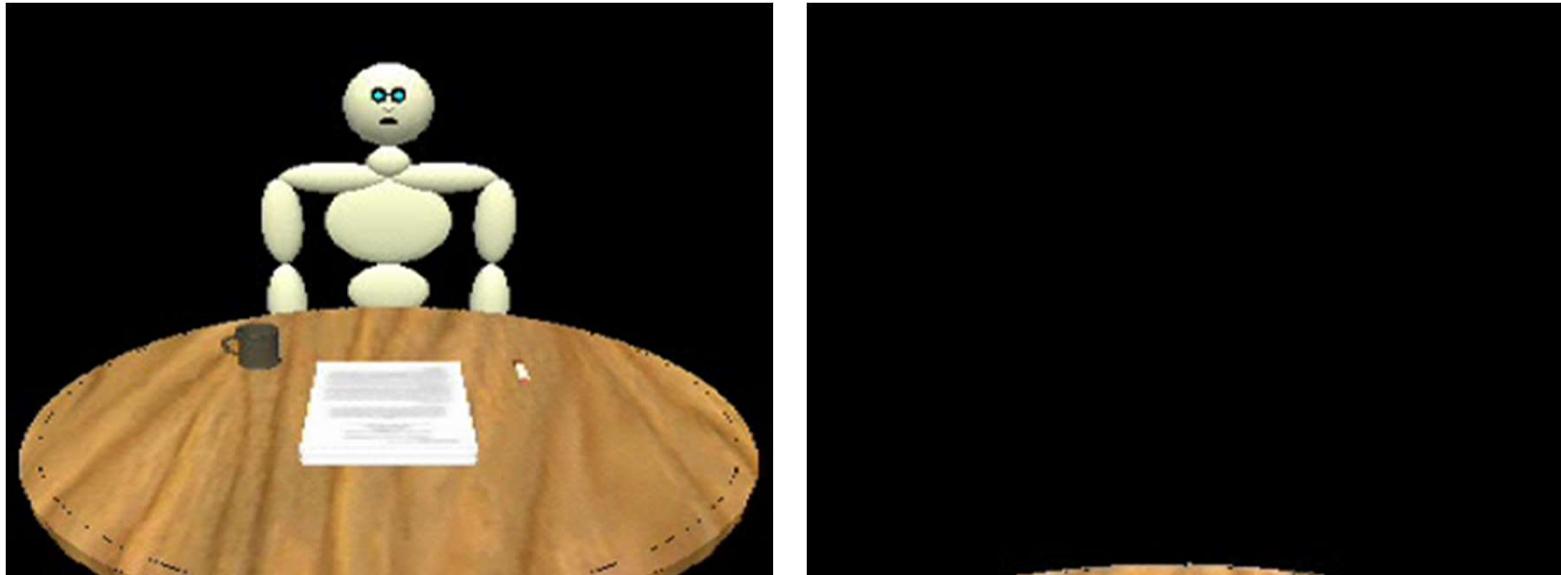
- ◆ Hierarchical and parallel state machines
 - Hierarchy = filtering functions
 - Hierarchy = composition of parallel state machines



Finite state machines

- ◆ Can be used for behavior and scenarios descriptions
- ◆ Pros
 - Simplicity of description
 - Consistency ensured by state machine structure
 - Perfect control
 - Possibility to add time information etc...
- ◆ Cons
 - Difficulty to avoid dead locks in case of parallel state machines
 - Hard to maintain when FSM become very complex

Finite state machines



Automatic scheduling of parallel FSM
Dead lock avoidance, priorities, preferences

F Lamarche, and S Donikian (2002). *Automatic Orchestration of Behaviours through the management of Resources and Priority Levels*. In: Autonomous Agents and Multi Agent Systems. ACM, Bologna, Italy

FSM and task modelling

- ◆ State machines

- Useful to describe atomic behaviours
- Needs a higher description level to describe complex behaviours

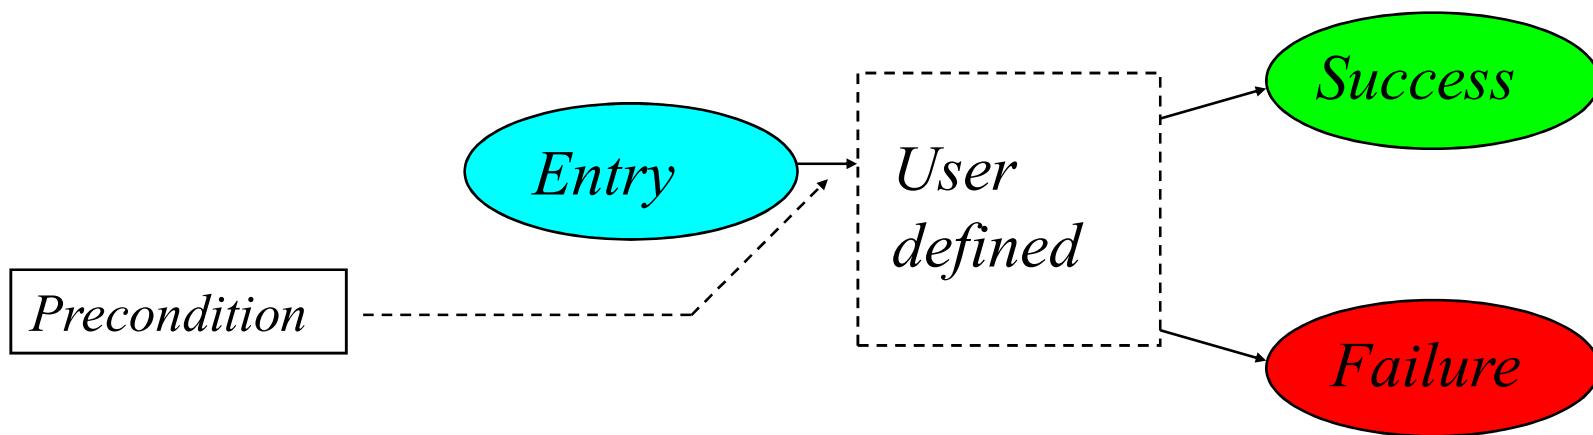
- ◆ Tasks

- Define a modelling framework
 - » Provide an abstract behaviour with identified states
- Uses inheritance and polymorphism
 - » behaviour implementation
- Operator definition
 - » High level task creation and manipulation

FSM and task modelling

◆ Task model

- Three predefined states : Entry / Success / Failure
- Precondition(s)
- Conserved resources
- User defined behaviour through inheritance



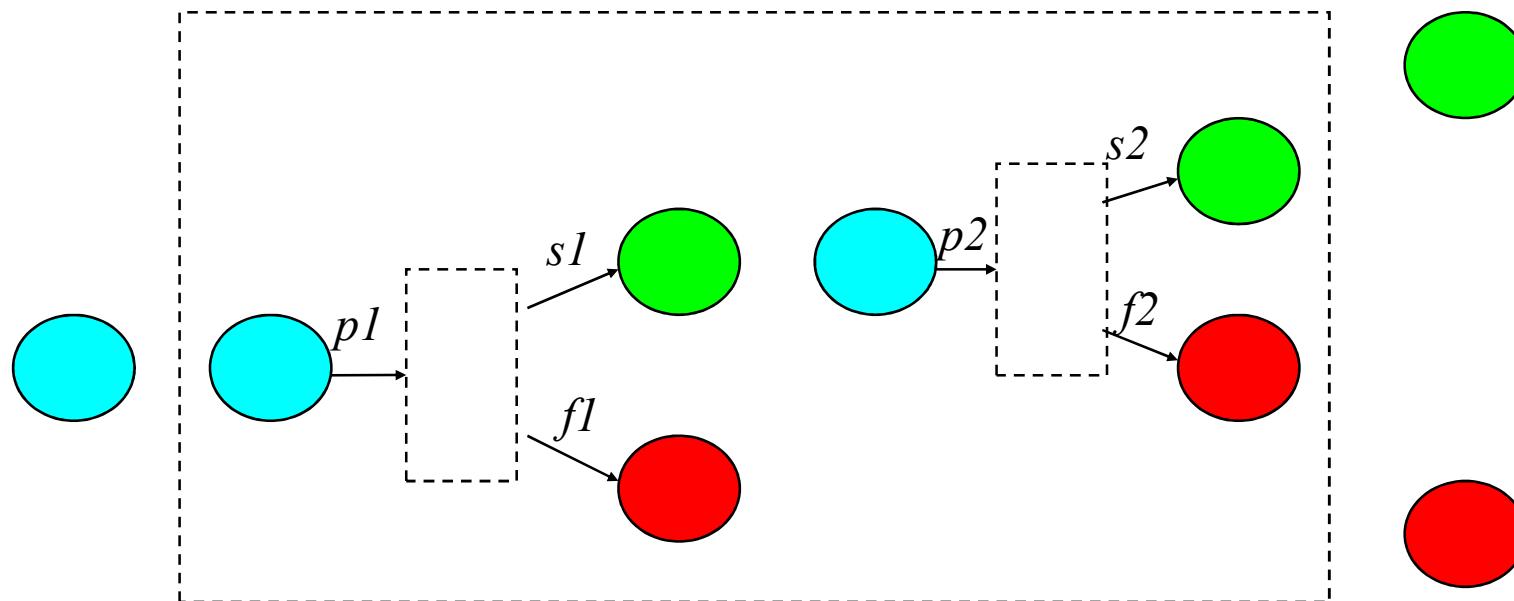
FSM and task modelling

- ◆ Task model provides operators
 - Logical combination of different tasks
 - Sequence / on failure / choice / all without order / parallel / repeat until / realize X among all
 - Parameters : tasks to combine
 - State machine generation with structure modification
- ◆ New task generation
 - Can be combined with other tasks thanks to operators

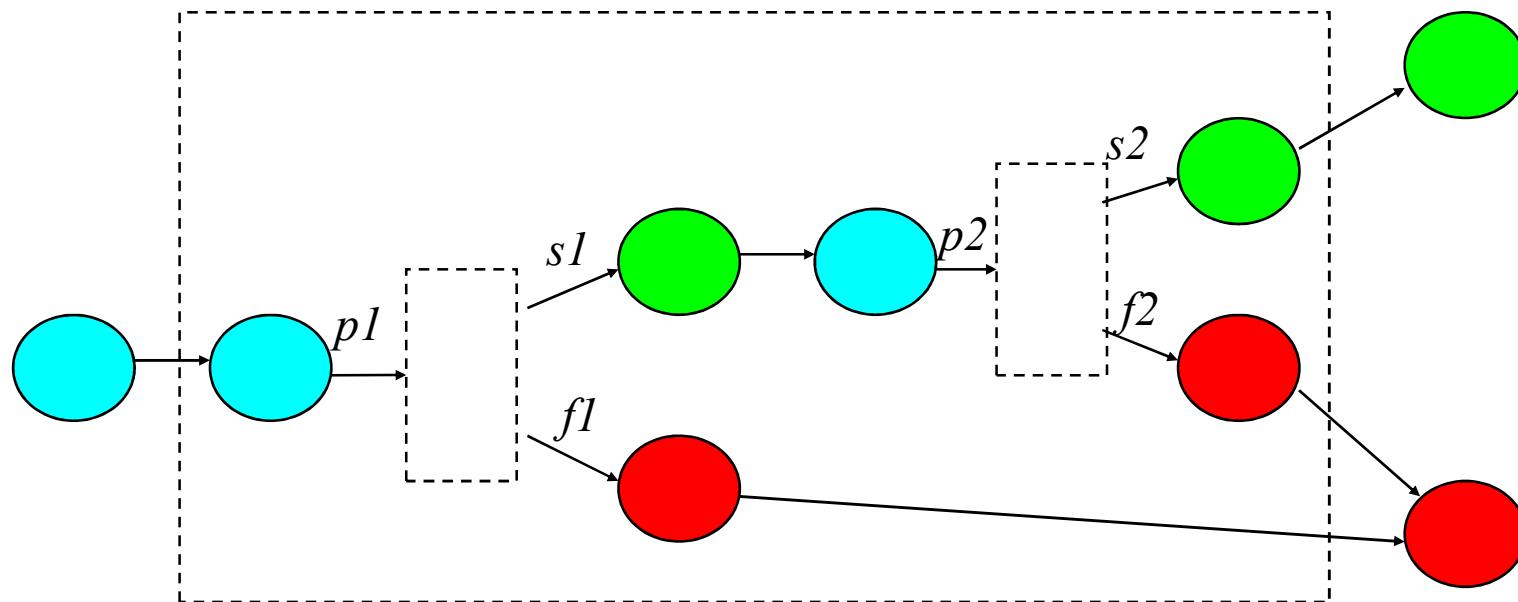
Example : Sequence operator



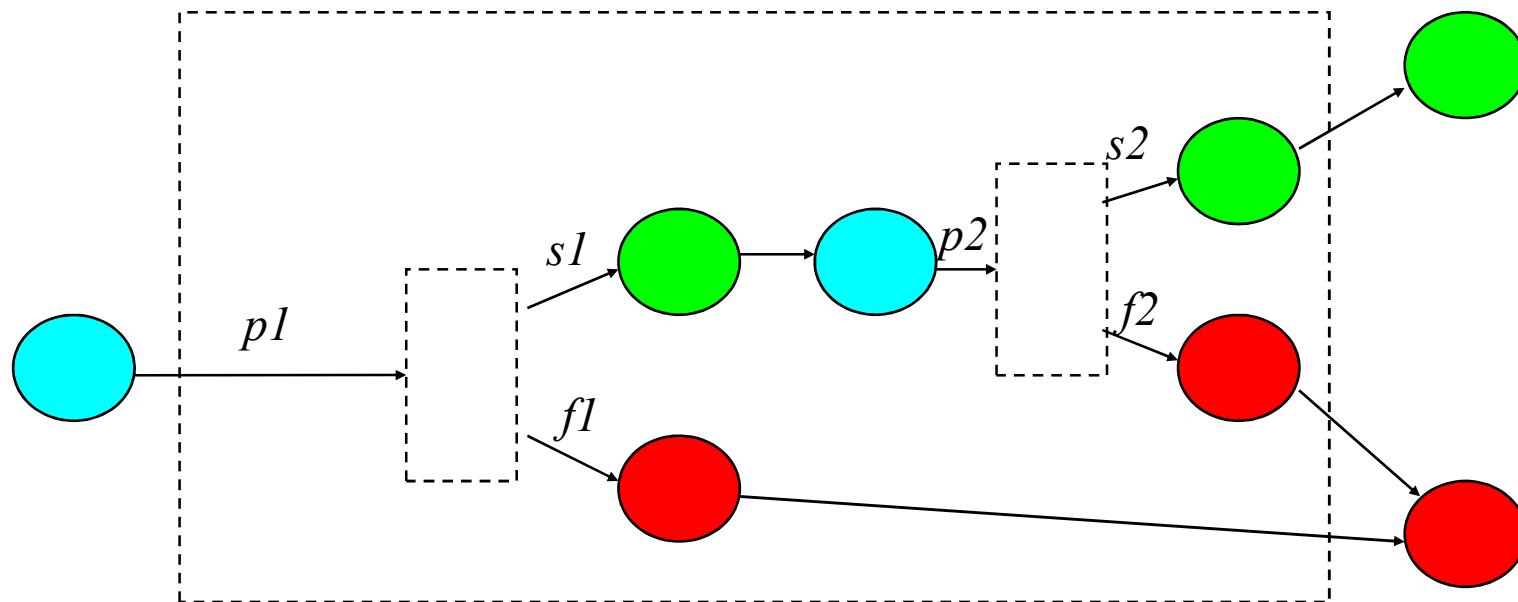
Example : Sequence operator



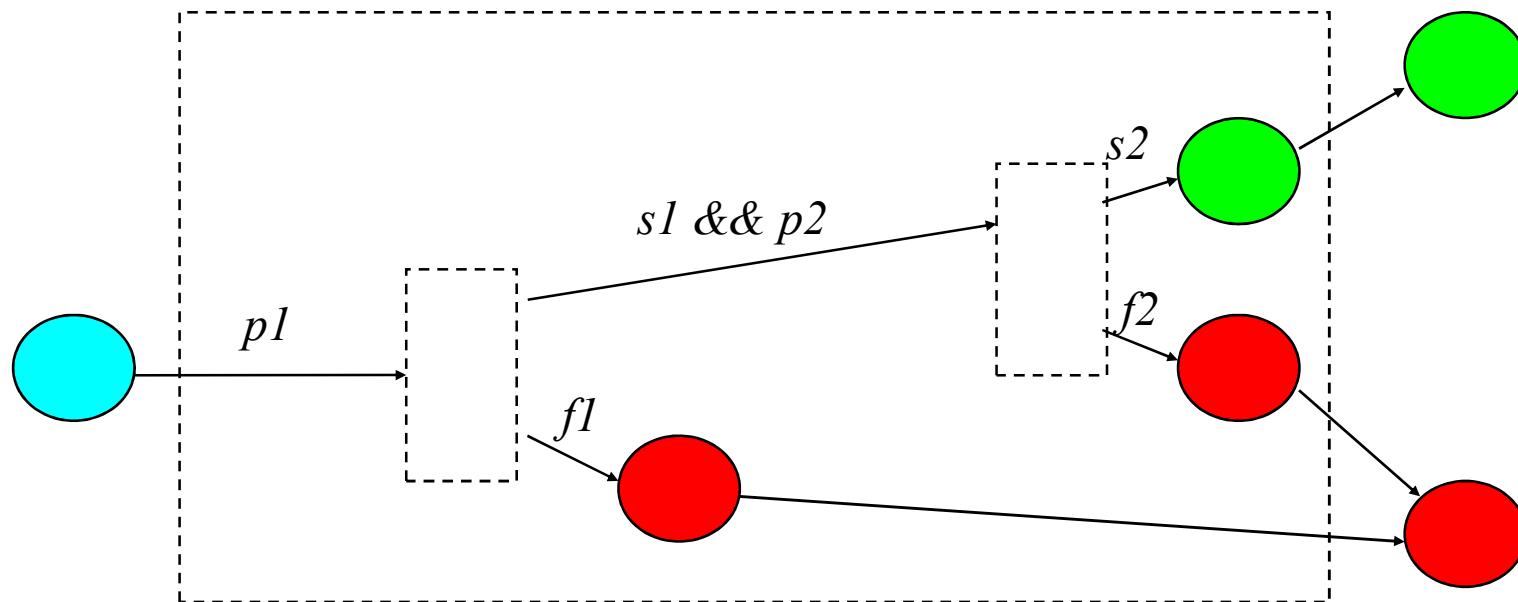
Example : Sequence operator



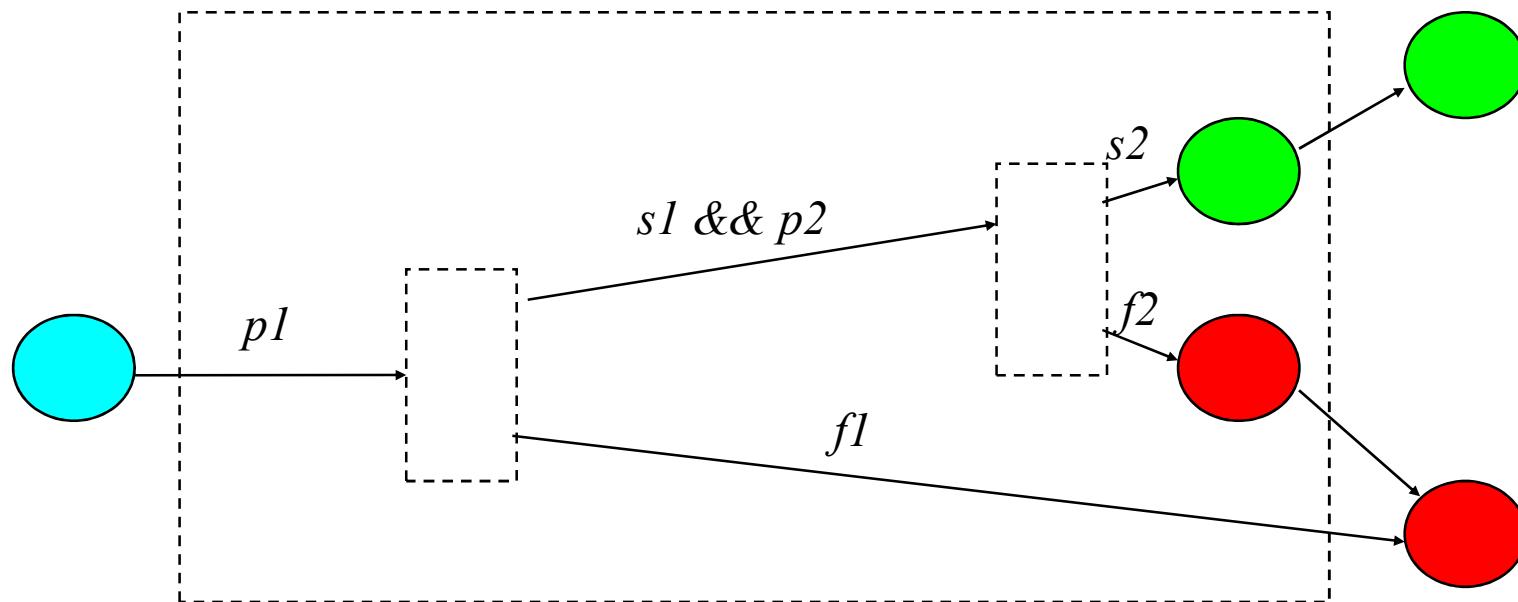
Example : Sequence operator



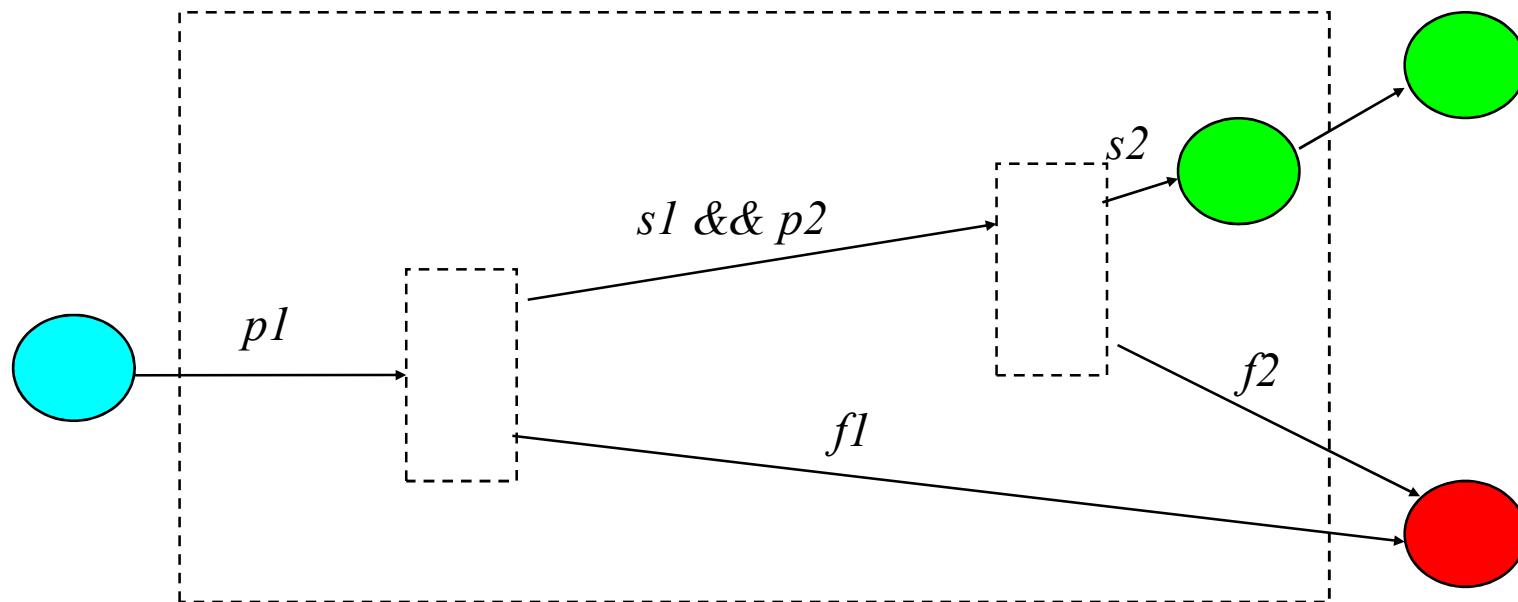
Example : Sequence operator



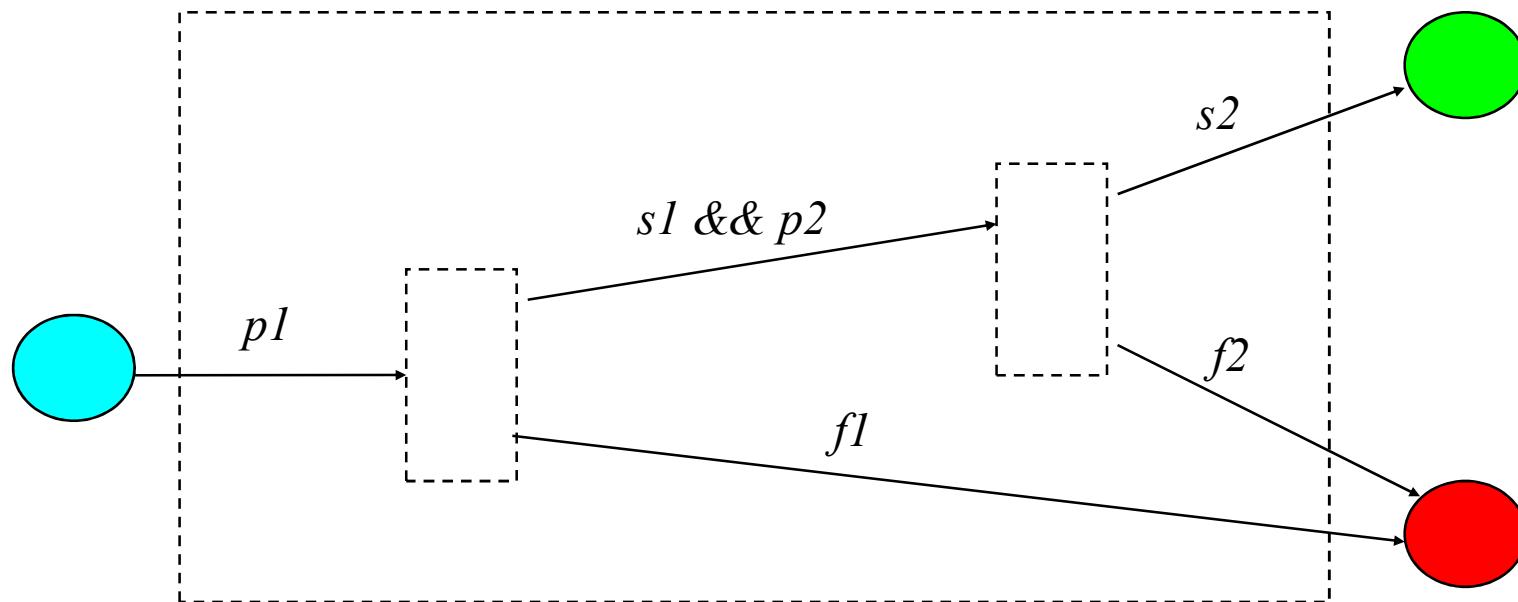
Example : Sequence operator



Example : Sequence operator



Example : Sequence operator



New task generation

FSM and task modeling

- ◆ FSM & Task modeling

- Generic framework providing high level operators
- Simple description of complex behaviours
- Easier to maintain than very complex hand designed state machines
- Reuse of pre-described state machines
 - » Modifications are automatically taken into account

- ◆ Remark

- Keep operators, encode them and replace FSMs by scripts / code and...
 - » **You obtain behavior trees**

Hierarchical task networks

- ◆ Classical planning:
 - Each state of the world is represented by a set of atoms
 - » Can be true or false / present or absent
 - » Relate to object state or relations between objects
 - » Ex: (closed door), (on-table my-computer), (at-location me classroom)
 - Each action (called operator) corresponds to a deterministic state transition
 - » Parameters: objects used in preconditions and effects
 - » Precondition: conjunction of atoms that must be present to execute the action
 - » Effect: adds atoms (add list) and remove some (del list)

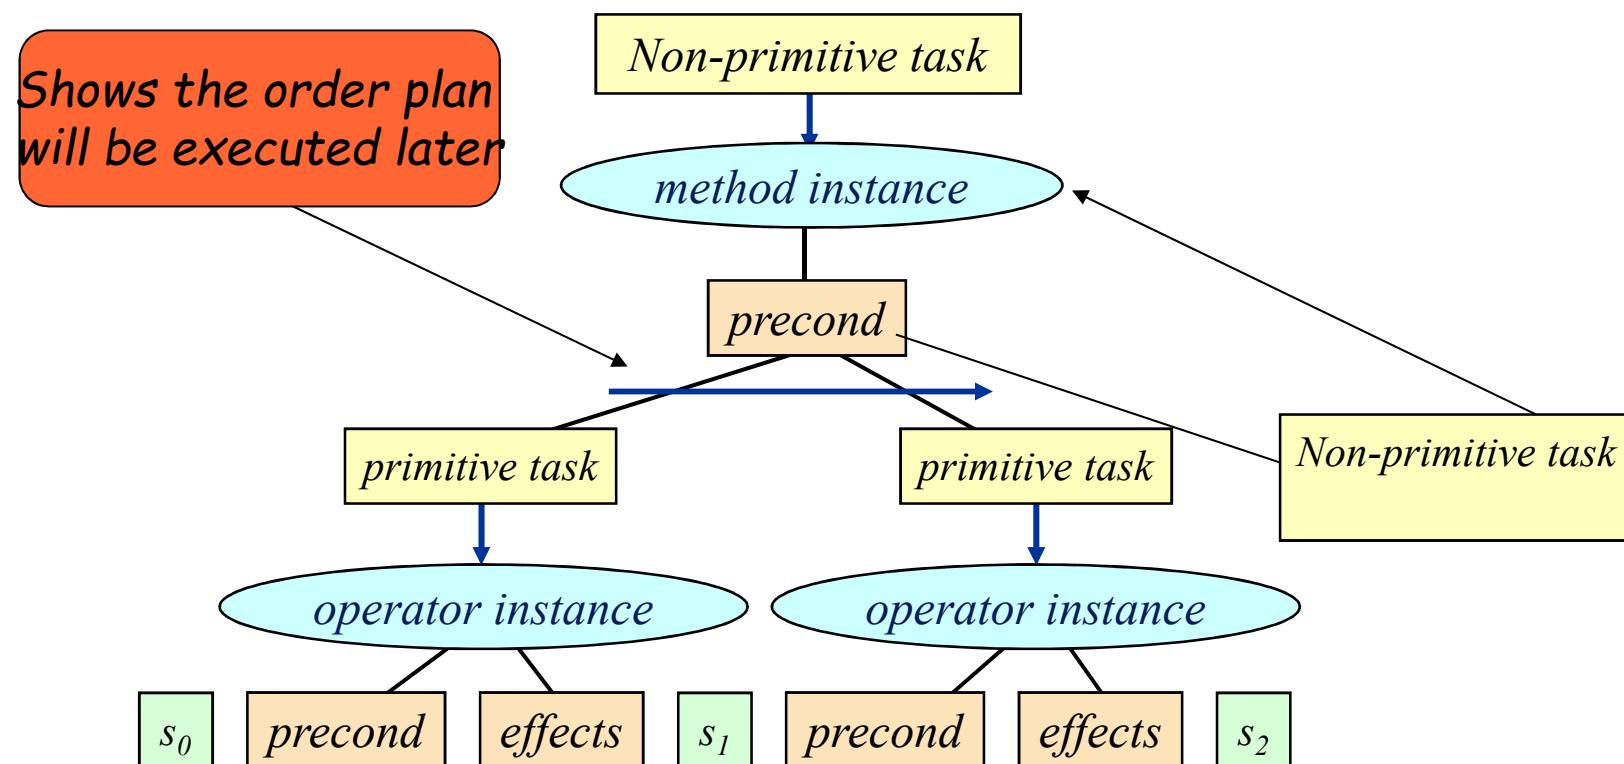
Hierarchical task networks

- ◆ Classical planning:
 - Goal : a sub state of the world
 - Ex : (in-corridor me)
 - Search for a succession of actions satisfying the goal

- ◆ But
 - We can't describe activities that cannot be expressed as a sub state of the world
 - » Ex: go for a walk...

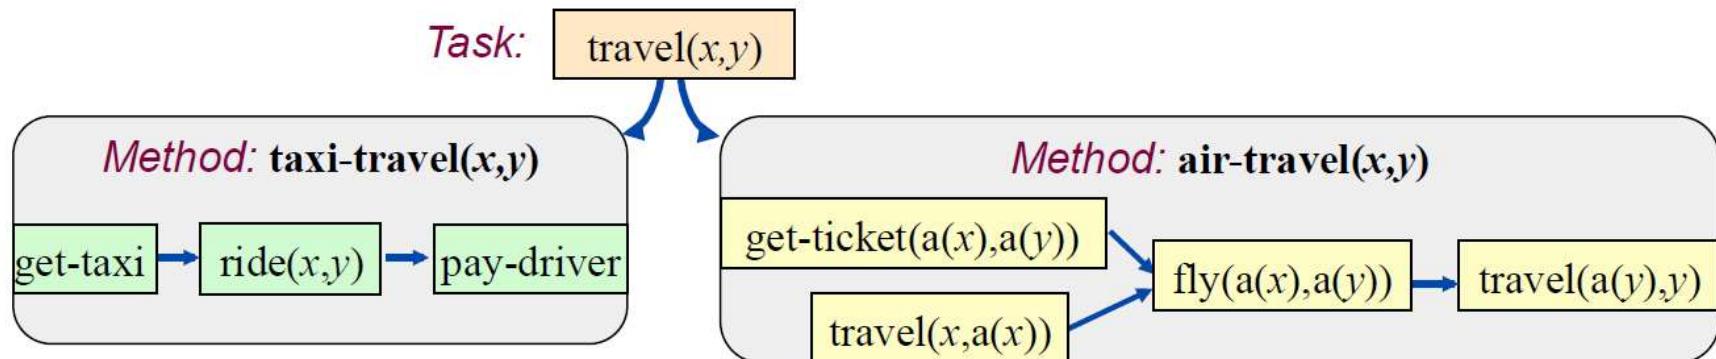
Hierarchical task networks

- ◆ Objective: to perform a set of tasks not a set of goals
- ◆ Adds notions of tasks, methods, task networks
- ◆ Tasks decompose into subtasks thanks to methods



Hierarchical task networks

- ◆ Example



Hierarchical task networks

- ◆ Method description

- Non primitive (decomposition into subtasks)
 - » Identifier and parameters
 - » Task performed by the method (with parameters)
 - » A precondition
 - » A sequence of subtasks

- ◆ Ex:

- Method: airtravel(x,y)
- Task: travel(x,y)
- Precondition: longdistance(x,y)
- Subtasks: buy-ticket(a(x), a(y)), travel(x,a(x)), fly(a(x), a(y)), travel(a(y),y)

Hierarchical task networks

- ◆ Method description
 - Primitive
 - » Identifier and parameters
 - » Task realized by the method (with parameters)
 - » Operator (with parameters)
- ◆ Relationship between non-primitive tasks and methods
 - Decomposed by applying a method
- ◆ Relationship between primitive tasks and operators
 - Primitive task is achieved by applying an operator

Hierarchical task networks

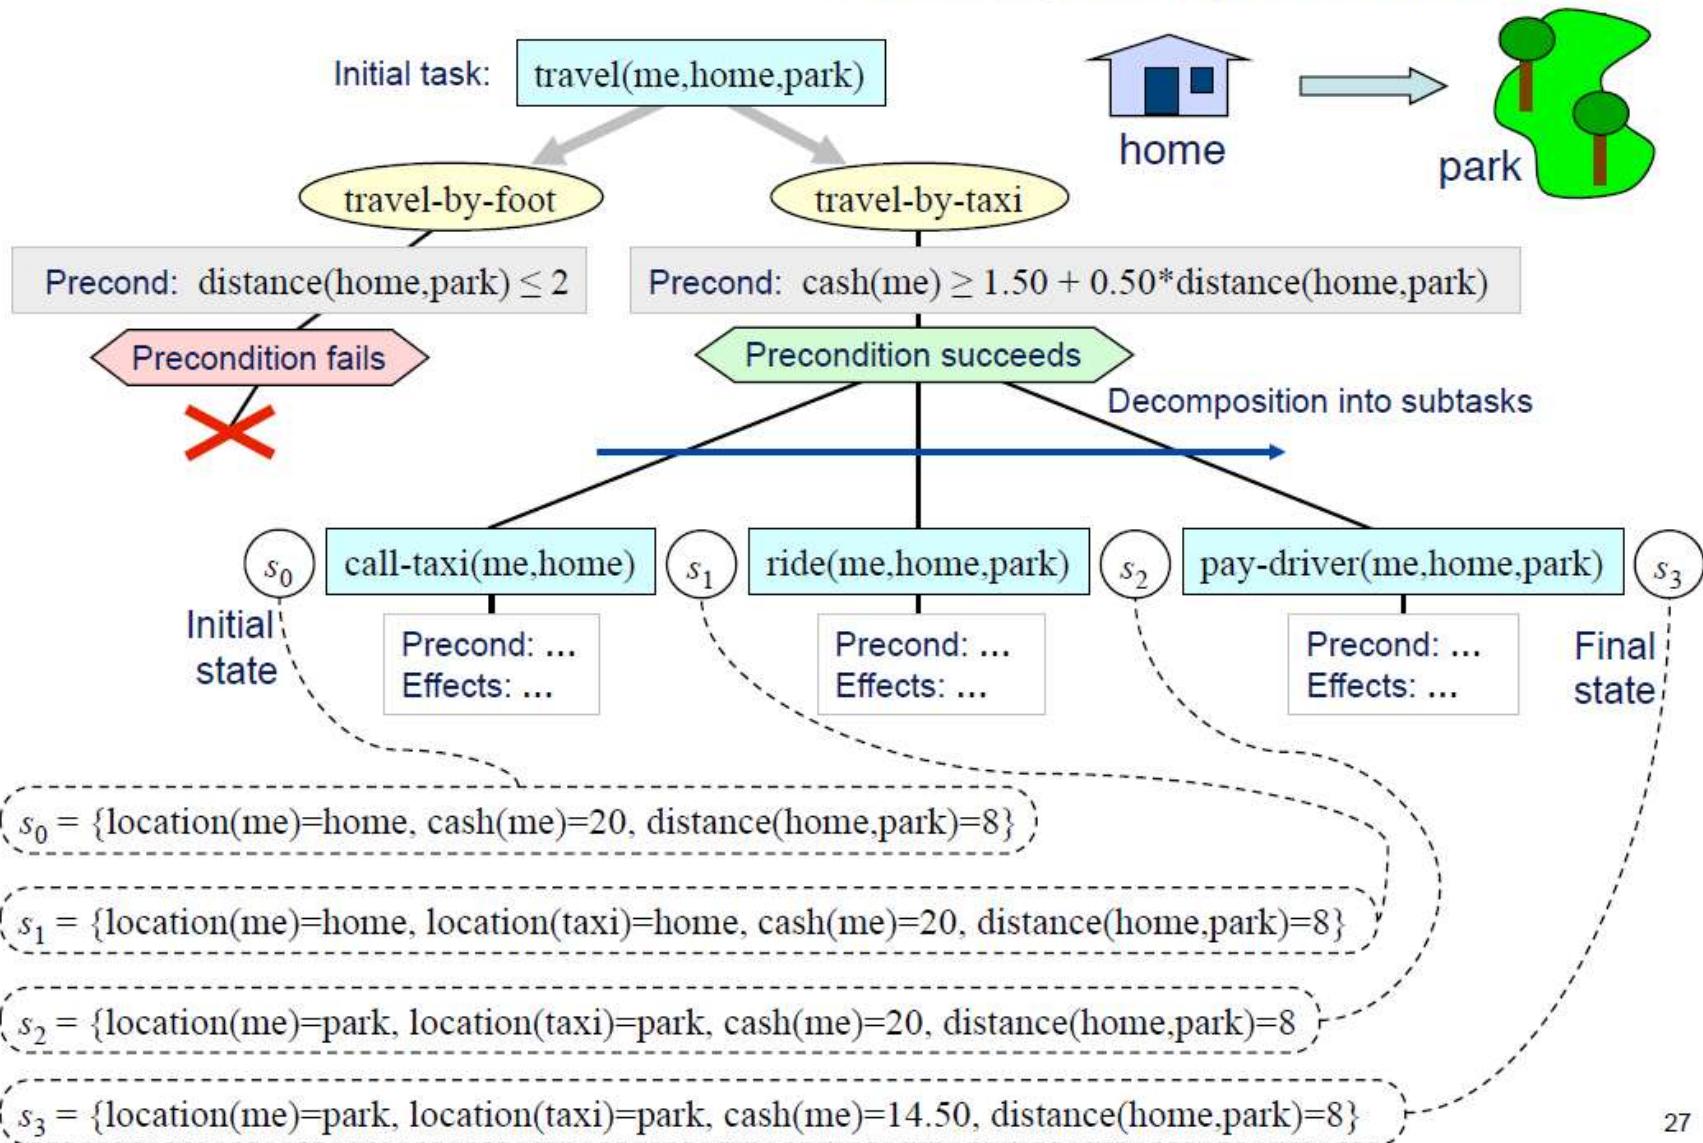
- ◆ Domain consists of
 - methods and operators (SHOP- axioms)
- ◆ Problem consists of
 - domain, initial state, initial task network (tasks to accomplish, with some ordering of the tasks defined)
- ◆ Solution
 - A plan: Totally ordered collection of primitive tasks

Hierarchical task networks

- ◆ Planning technique overview
 - Given a list of tasks
 - Use a left first exploration of the decomposition tree
 - When a free variable appears (often the case)
 - » Try multiple instantiations of this variable with objects of the world
 - When a precondition is not satisfied backtrack either by
 - » Trying another free variable instantiation if possible
 - » Returning a failure and trying another decomposition
 - When an operator appear an its precondition is true
 - » Apply the operator to change the state of the world
 - » Continue decomposition tree exploration
- ◆ At the end, either a valid decomposition has been found (a plan has been found) or the planner failed.

Planning Problem:

I am at home, I have \$20,
I want to go to a park 8 miles away



Hierarchical task networks

- ◆ Some implementations exists and can be downloaded
 - SHOP, JSHOP, UMCP
- ◆ Mixes long term planning and designer control by method description
- ◆ More expressive than classical planning
 - Methods, tasks instead of goals
 - Numerical values in preconditions and effects (ressources)
- ◆ But not reactive...